

MNIST-Pytorch

January 24, 2021

0.0.1 Convolutional Neural Network

```
[1]: import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data import DataLoader
from torchvision import datasets, transforms
from torchvision.utils import make_grid

import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, classification_report
plt.style.use('seaborn-white')
import warnings
warnings.filterwarnings('ignore', 'DeprecatedWarnings')
```

```
[2]: torch.__version__
```

```
[2]: '1.7.1'
```

0.0.2 Loading Datasets

```
[3]: # MNIST Is considered to be the Hello World of CNN's and Deep Learning in
    ↪General, so it is
# generally provided from as the inbuild go-to project for beginners
transforms = transforms.ToTensor()
train_data = datasets.MNIST(root='../Data', train=True, download=True,
    ↪transform=transforms)
test_data = datasets.MNIST(root='../Data', train=False, download=True,
    ↪transform=transforms)
```

```
[4]: train_data
```

```
[4]: Dataset MNIST
      Number of datapoints: 60000
      Root location: ../Data
      Split: Train
      StandardTransform
```

```
Transform: ToTensor()
```

```
[5]: test_data
```

```
[5]: Dataset MNIST
      Number of datapoints: 10000
      Root location: ../Data
      Split: Test
      StandardTransform
      Transform: ToTensor()
```

0.0.3 Data Loader Object

```
[6]: # Pytorch's convention to create dataloader objects, to be used in the models
      ↪train cycle
      batch_s = 10
      train_loader = DataLoader(train_data, batch_size=batch_s, shuffle=True)
      test_loader = DataLoader(test_data, batch_size=batch_s, shuffle=False)
```

0.0.4 What is the Data about?

```
[7]: # pipeline obj=1/
      desc = '''
      The MNIST database of handwritten digits, available from this page, has a
      ↪training set of 60,000 examples, and a test set of 10,000 examples. It is a
      ↪subset of a larger set available from NIST. The digits have been
      ↪size-normalized and centered in a fixed-size image.
      '''
      print(desc)
```

The MNIST database of handwritten digits, available from this page, has a training set of 60,000 examples, and a test set of 10,000 examples. It is a subset of a larger set available from NIST. The digits have been size-normalized and centered in a fixed-size image.


```
[8]: (x_train, y_train) = train_data[0] # one sample
```

```
[9]: y_train # label
```

```
[9]: 5
```

```
[10]: plt.imshow(x_train[0], cmap='binary')
      plt.title('label {}'.format(y_train))
```

```
[10]: Text(0.5, 1.0, 'label 5')
```



0.0.5 Convolutional Neural Network Class

0.0.6 Buidling Model

```
[11]: class convolutional(nn.Module):

    def __init__(self,in_channel=1,output=10):
        super().__init__() # inheriting nn.Module
        # defining Layer Functions
        # input dimension = (b,28,28,1)
        self.conv1 = nn.Conv2d(in_channels=in_channel, out_channels=6,
        ↪kernel_size=(3,3), stride=1 )
        # (b,26,26,6)
        self.pool = nn.MaxPool2d(kernel_size=(2,2), stride=(2,2)) # f=2, s=2
        # dimension (b,12,12,6)
        self.conv2 = nn.
        ↪Conv2d(in_channels=6,out_channels=16,kernel_size=(3,3),stride=1)
        # dimension (b,10,10,16)
        # we apply one more pool after this dimension so (b,5,5,16)
        self.fc1 = nn.Linear(5*5*16,120)
        self.fc2 = nn.Linear(120, 64)
        self.fc3 = nn.Linear(in_features=64,out_features=output)
```

```

def forward(self,X):
    # Construction of the Architecture
    X = F.relu(self.conv1(X))
    X = self.pool(X)
    X = F.relu(self.conv2(X))
    X = self.pool(X)
    X = X.view(-1, 5*5*16) # Flatten
    X = F.relu(self.fc1(X))
    X = F.relu(self.fc2(X))
    X = self.fc3(X)
    return F.log_softmax(X, dim=1)

```

```
[12]: net = convolutional()
```

```
[13]: net
```

```

[13]: convolutional(
      (conv1): Conv2d(1, 6, kernel_size=(3, 3), stride=(1, 1))
      (pool): MaxPool2d(kernel_size=(2, 2), stride=(2, 2), padding=0, dilation=1,
      ceil_mode=False)
      (conv2): Conv2d(6, 16, kernel_size=(3, 3), stride=(1, 1))
      (fc1): Linear(in_features=400, out_features=120, bias=True)
      (fc2): Linear(in_features=120, out_features=64, bias=True)
      (fc3): Linear(in_features=64, out_features=10, bias=True)
)

```

```

[14]: print('Total Learnable Params :')
      params = [p.numel() for p in net.parameters() if p.requires_grad]
      for item in params:
          print(f'{item:>10}')
      print(f'-----\n{sum(params):>10}')

```

Total Learnable Params :

```

      54
       6
      864
       16
     48000
       120
      7680
        64
       640
        10

```

57454

```
[15]: # defining the loss function and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(net.parameters(),lr=0.001)

[17]: # The TRAIN CYCLE
import time
start = time.perf_counter()
# trackers, for post evaluation
epochs = 5
train_losses = []
test_losses = []
train_correct = []
test_correct = []
accuracy = []

for i in range(epochs):
    trn_corr = 0
    tst_corr = 0

    # Run the training batches
    for b, (X_train, y_train) in enumerate(train_loader):
        b+=1

        # Apply the model
        y_pred = net(X_train) # we don't flatten X-train here
        loss = criterion(y_pred, y_train)

        # Tally the number of correct predictions
        _, predicted = y_pred.max(1)
        batch_corr = (predicted == y_train).sum()
        trn_corr += batch_corr

        # Update parameters
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        # Print interim results
        if b%600 == 0:
            print(f'epoch: {i:2}  batch: {b:4} [{10*b:6}/60000]  loss: {loss.
↪item():10.8f}  \
accuracy: {trn_corr.item()*100/(10*b):7.3f}%' )
            accuracy.append(trn_corr.item()*100/(10*b))
```

```

train_losses.append(loss)
train_correct.append(trn_corr)

# Run the testing batches
with torch.no_grad():
    for b, (X_test, y_test) in enumerate(test_loader):

        # Apply the model
        y_val = net(X_test)

        # Tally the number of correct predictions
        predicted = torch.max(y_val.data, 1)[1]
        tst_corr += (predicted == y_test).sum()

    loss = criterion(y_val, y_test)
    test_losses.append(loss)
    test_correct.append(tst_corr)

elapsed = (time.perf_counter() - start)/60
print('ELAPSED {:.2f}'.format(elapsed))

```

epoch:	0	batch:	600	[6000/60000]	loss:	0.11924569	accuracy:	75.550%
epoch:	0	batch:	1200	[12000/60000]	loss:	0.12763628	accuracy:	84.292%
epoch:	0	batch:	1800	[18000/60000]	loss:	0.04454534	accuracy:	87.811%
epoch:	0	batch:	2400	[24000/60000]	loss:	0.00629141	accuracy:	89.662%
epoch:	0	batch:	3000	[30000/60000]	loss:	0.05193086	accuracy:	90.840%
epoch:	0	batch:	3600	[36000/60000]	loss:	0.02835705	accuracy:	91.792%
epoch:	0	batch:	4200	[42000/60000]	loss:	0.56388974	accuracy:	92.493%
epoch:	0	batch:	4800	[48000/60000]	loss:	0.00274723	accuracy:	93.071%
epoch:	0	batch:	5400	[54000/60000]	loss:	0.08888493	accuracy:	93.496%
epoch:	0	batch:	6000	[60000/60000]	loss:	0.00419691	accuracy:	93.852%
epoch:	1	batch:	600	[6000/60000]	loss:	0.00491270	accuracy:	97.450%
epoch:	1	batch:	1200	[12000/60000]	loss:	0.00427231	accuracy:	97.458%
epoch:	1	batch:	1800	[18000/60000]	loss:	0.04439243	accuracy:	97.544%
epoch:	1	batch:	2400	[24000/60000]	loss:	0.01586340	accuracy:	97.588%
epoch:	1	batch:	3000	[30000/60000]	loss:	0.02111457	accuracy:	97.633%
epoch:	1	batch:	3600	[36000/60000]	loss:	0.00026605	accuracy:	97.717%
epoch:	1	batch:	4200	[42000/60000]	loss:	0.26243860	accuracy:	97.743%
epoch:	1	batch:	4800	[48000/60000]	loss:	0.29612529	accuracy:	97.783%
epoch:	1	batch:	5400	[54000/60000]	loss:	0.13613828	accuracy:	97.850%
epoch:	1	batch:	6000	[60000/60000]	loss:	0.16438568	accuracy:	97.848%
epoch:	2	batch:	600	[6000/60000]	loss:	0.00097367	accuracy:	98.367%
epoch:	2	batch:	1200	[12000/60000]	loss:	0.00730920	accuracy:	98.342%
epoch:	2	batch:	1800	[18000/60000]	loss:	0.44688016	accuracy:	98.356%
epoch:	2	batch:	2400	[24000/60000]	loss:	0.02830379	accuracy:	98.350%
epoch:	2	batch:	3000	[30000/60000]	loss:	0.16491416	accuracy:	98.333%

```

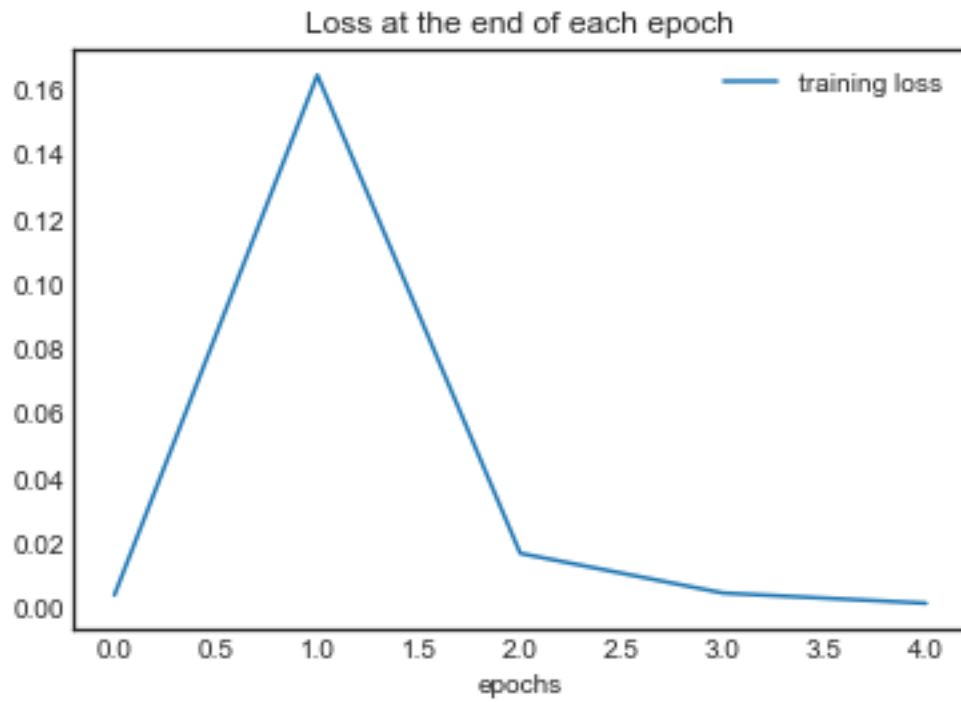
epoch: 2 batch: 3600 [ 36000/60000] loss: 0.31121570 accuracy: 98.328%
epoch: 2 batch: 4200 [ 42000/60000] loss: 0.83973742 accuracy: 98.360%
epoch: 2 batch: 4800 [ 48000/60000] loss: 0.02205916 accuracy: 98.369%
epoch: 2 batch: 5400 [ 54000/60000] loss: 0.01329469 accuracy: 98.387%
epoch: 2 batch: 6000 [ 60000/60000] loss: 0.01703760 accuracy: 98.410%
epoch: 3 batch: 600 [ 6000/60000] loss: 0.00008939 accuracy: 98.950%
epoch: 3 batch: 1200 [ 12000/60000] loss: 0.00989255 accuracy: 98.833%
epoch: 3 batch: 1800 [ 18000/60000] loss: 0.01814078 accuracy: 98.850%
epoch: 3 batch: 2400 [ 24000/60000] loss: 0.00006347 accuracy: 98.867%
epoch: 3 batch: 3000 [ 30000/60000] loss: 0.08430563 accuracy: 98.833%
epoch: 3 batch: 3600 [ 36000/60000] loss: 0.00083146 accuracy: 98.814%
epoch: 3 batch: 4200 [ 42000/60000] loss: 0.00024983 accuracy: 98.790%
epoch: 3 batch: 4800 [ 48000/60000] loss: 0.04406525 accuracy: 98.785%
epoch: 3 batch: 5400 [ 54000/60000] loss: 0.00120888 accuracy: 98.794%
epoch: 3 batch: 6000 [ 60000/60000] loss: 0.00479843 accuracy: 98.792%
epoch: 4 batch: 600 [ 6000/60000] loss: 0.03544014 accuracy: 98.967%
epoch: 4 batch: 1200 [ 12000/60000] loss: 0.00299220 accuracy: 99.017%
epoch: 4 batch: 1800 [ 18000/60000] loss: 0.02138555 accuracy: 99.000%
epoch: 4 batch: 2400 [ 24000/60000] loss: 0.00595939 accuracy: 98.971%
epoch: 4 batch: 3000 [ 30000/60000] loss: 0.00030641 accuracy: 98.980%
epoch: 4 batch: 3600 [ 36000/60000] loss: 0.00556433 accuracy: 99.000%
epoch: 4 batch: 4200 [ 42000/60000] loss: 0.07285134 accuracy: 98.936%
epoch: 4 batch: 4800 [ 48000/60000] loss: 0.00090826 accuracy: 98.973%
epoch: 4 batch: 5400 [ 54000/60000] loss: 0.00183687 accuracy: 98.954%
epoch: 4 batch: 6000 [ 60000/60000] loss: 0.00166669 accuracy: 98.988%
ELAPSED 2.48

```

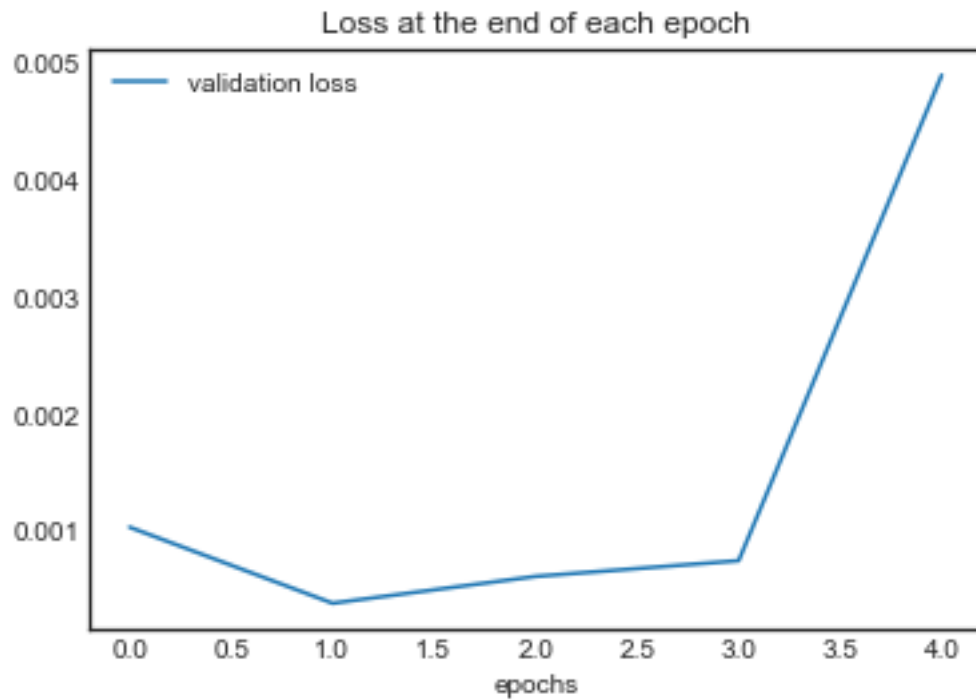
```

[18]: plt.plot(train_losses, label='training loss')
      plt.title('Loss at the end of each epoch')
      plt.xlabel('epochs')
      plt.legend();


```




```
[19]: plt.plot(test_losses, label='validation loss')
plt.title('Loss at the end of each epoch')
plt.xlabel('epochs')
plt.legend();
```

```
[20]: plt.plot([t/600 for t in train_correct], label='training accuracy')  
  
plt.title('Accuracy at the end of each epoch')  
plt.legend();
```



```
[21]: plt.plot([t/100 for t in test_correct], label='validation accuracy')
plt.title('Accuracy at the end of each epoch')
plt.legend();
```



0.0.7 Overall Evaluation

```
[22]: test_loader = DataLoader(test_data, batch_size=10000, shuffle=False)
```

```
[23]: with torch.no_grad():
    correct = 0
    for x_test, y_test in test_loader:
        y_val = net(x_test)
        _, predicted = y_val.max(1)
        correct += (predicted == y_test).sum()

print(f'Test accuracy: {correct.item()}/{len(test_data)} = {correct.item()*100/
    ↳ (len(test_data)):7.3f}%')
```

Test accuracy: 9872/10000 = 98.720%

```
[24]: predicted
```

```
[24]: tensor([7, 2, 1, ..., 4, 5, 6])
```

```
[25]: # Missed Pictures and their indexes
misses = np.array([])

for i in range(len(predicted.view(-1))):

    if predicted[i] != y_test[i]:
        misses = np.append(misses,i).astype('int64')


misses
```

```
[25]: array([ 18,  62, 115, 158, 340, 359, 551, 582, 619, 629, 659,
          684, 716, 720, 740, 924, 947, 1014, 1039, 1112, 1226, 1242,
          1247, 1299, 1328, 1393, 1500, 1522, 1621, 1709, 1721, 1737, 1782,
          1790, 1878, 1901, 1909, 2018, 2024, 2035, 2070, 2090, 2130, 2135,
          2293, 2343, 2447, 2488, 2496, 2582, 2597, 2648, 2654, 2743, 2771,
          2927, 2939, 2953, 2959, 2970, 3030, 3225, 3289, 3448, 3520, 3702,
          3726, 3756, 3762, 3767, 3778, 3806, 3808, 3902, 3906, 3941, 4163,
          4176, 4194, 4201, 4284, 4360, 4497, 4504, 4536, 4571, 4740, 4783,
          4838, 4860, 4880, 4911, 5127, 5457, 5634, 5642, 5937, 5955, 5973,
          5981, 5997, 6101, 6560, 6571, 6576, 6597, 6625, 6651, 6783, 6883,
          7216, 7574, 7899, 7928, 7990, 8279, 8325, 8469, 8509, 8520, 9009,
          9540, 9669, 9729, 9839, 9891, 9904, 9905])
```

0.0.8 Single Image Test

```
[26]: index = 2020
plt.figure(figsize=(4,4))
plt.imshow(x_test[index].reshape(28,28),cmap='binary')
plt.title('Label {}'.format(y_test[2020]))
```

```
[26]: Text(0.5, 1.0, 'Label 3')
```



```
[27]: net.eval()
      with torch.no_grad():
          yhat = net(x_test[index][0].view(1,1,28,28))
          print('Prediction ~ {}'.format(yhat.argmax()))
          print('Reality - {}'.format(y_test[index]))
```

Prediction ~ 3

Reality - 3

0.0.9 Miscellaneous

```
[28]: # Pipeline Script Data
```

```
[89]: project_name = 'MNIST'
      framework = 'Pytorch'
      prediction_type = 'Classification'
      network_type = 'Convolutional Neural Network'
      Architecture = str(net)
      layers = 6
      hidden_units = None
      Activations = "['relu','softmax']"
      epochs = 5
      metrics = 'Accuracy'
      Train_Accuracy = np.mean(accuracy)
      Test_Accuracy = 98.720
```

```
elapsed = '2.48 Min'
Desc = desc.strip()
```

```
[90]: param = dict()
      var = ['project_name', 'framework', 'prediction_type', 'network_type',
            'Architecture', 'layers', 'hidden_units', 'Activations', 'epochs',
            'metrics', 'Train_Accuracy', 'Test_Accuracy', 'elapsed', 'Desc']
```

```
[91]: for val in var:
      param[val] = eval(val)
```

```
[92]: import pickle
      file = open("state_dict.txt", "wb")

      dictionary = param

      pickle.dump(dictionary, file)
      file.close()
```

```
[95]: param
```

```
[95]: {'project_name': 'MNIST',
      'framework': 'Pytorch',
      'prediction_type': 'Classification',
      'network_type': 'Convolutional Neural Network',
      'Architecture': 'convolutional(\n (conv1): Conv2d(1, 6, kernel_size=(3, 3),
stride=(1, 1))\n (pool): MaxPool2d(kernel_size=(2, 2), stride=(2, 2),
padding=0, dilation=1, ceil_mode=False)\n (conv2): Conv2d(6, 16,
kernel_size=(3, 3), stride=(1, 1))\n (fc1): Linear(in_features=400,
out_features=120, bias=True)\n (fc2): Linear(in_features=120, out_features=64,
bias=True)\n (fc3): Linear(in_features=64, out_features=10, bias=True)\n)',
      'layers': 6,
      'hidden_units': None,
      'Activations': "['relu', 'softmax']",
      'epochs': 5,
      'metrics': 'Accuracy',
      'Train_Accuracy': 96.62335542328043,
      'Test_Accuracy': 98.72,
      'elapsed': '2.48 Min',
      'Desc': 'The MNIST database of handwritten digits, available from this page,
has a training set of 60,000 examples, and a test set of 10,000 examples. It is
a subset of a larger set available from NIST. The digits have been size-
normalized and centered in a fixed-size image.'}
```

```
[99]: import pandas as pd
      pd.DataFrame(pd.Series(param)).T
```

```

[99]: project_name framework prediction_type network_type \
0      MNIST      Pytorch  Classification  Convolutional Neural Network

      Architecture layers hidden_units \
0  convolutional(\n (conv1): Conv2d(1, 6, kernel...      6      None

      Activations epochs  metrics Train_Accuracy Test_Accuracy  elapsed \
0  ['relu','softmax']      5  Accuracy      96.6234      98.72  2.48 Min

      Desc
0  The MNIST database of handwritten digits, avai...

```

```
[ ]:
```