

Dogs&Cats-classifier

February 23, 2021

0.0.1 Cats and Dogs Classifier using Convolutional Neural Networks

```
[1]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
[2]: ! ls drive/MyDrive/CATS_DOGS.zip
```

drive/MyDrive/CATS_DOGS.zip

```
[4]: # libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
plt.style.use('classic')
```

0.0.2 Paths

```
[5]: # archive = '../Data/CATS_DOGS.zip' # local
archive = 'drive/MyDrive/CATS_DOGS.zip' # colab
```

```
[6]: # unzipping
import zipfile
zip_obj = zipfile.ZipFile(file=archive)
```

```
[7]: import os
if os.path.exists('drive/MyDrive/cd_data/'):
    print(True)
else:
    print(False)
    print(os.mkdir('drive/MyDrive/cd_data/'))
    print('created!')
```

False

None

created!

```
[8]: # # extract to the location

# # zip_obj.extractall('../Data/') # local
zip_obj.extractall('drive/MyDrive/cd_data/') # colab
```

```
[9]: # train and validation paths

# train_dir = '../Data/CATS_DOGS/train/'
# valid_dir = '../Data/CATS_DOGS/test/'

train_dir = 'drive/MyDrive/cd_data/CATS_DOGS/train/'
valid_dir = 'drive/MyDrive/cd_data/CATS_DOGS/test/'
```

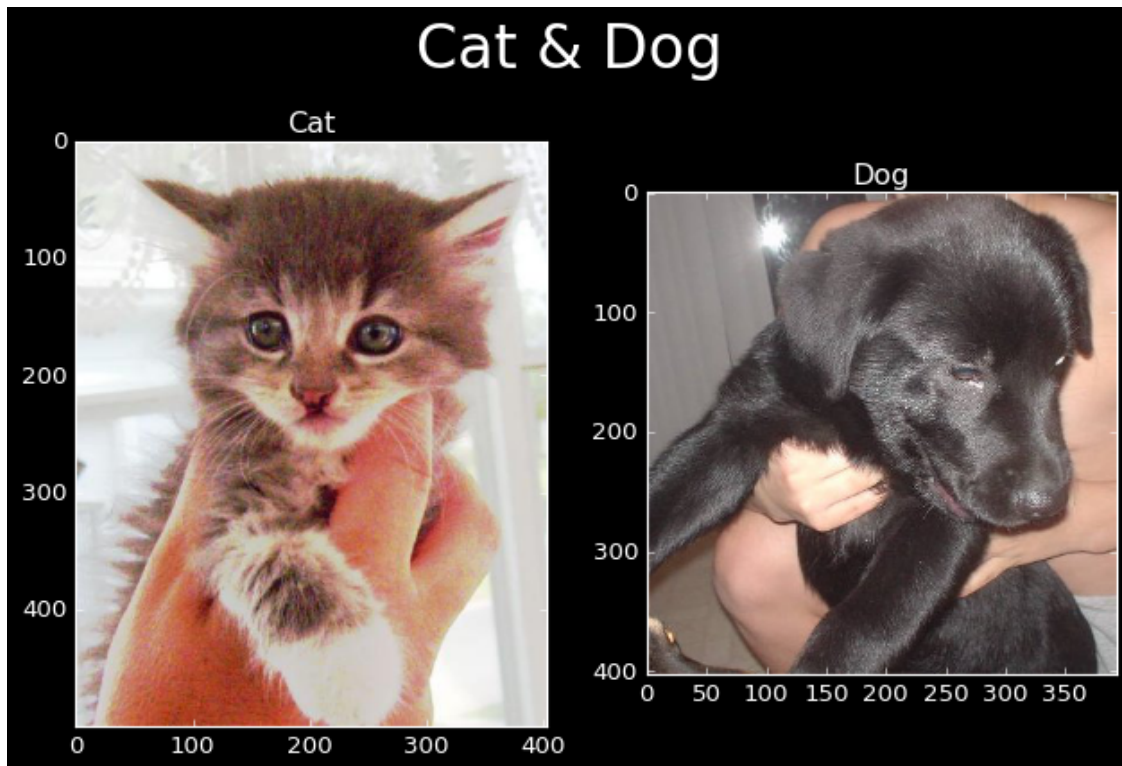
0.0.3 Raw Data

```
[10]: # random

# cat = '../Data/CATS_DOGS/train/CAT/100.jpg'
# dog = '../Data/CATS_DOGS/train/DOG/1050.jpg'

cat = 'drive/MyDrive/cd_data/CATS_DOGS/train/CAT/100.jpg'
dog = 'drive/MyDrive/cd_data/CATS_DOGS/train/DOG/1050.jpg'
```

```
[11]: plt.style.use('dark_background')
from PIL import Image
img_cat = Image.open(cat)
img_dog = Image.open(dog)
fig, axes = plt.subplots(1,2)
axes[0].imshow(np.array(img_cat),label='cat')
axes[0].set_title('Cat')
axes[1].imshow(np.array(img_dog),label='dog')
axes[1].set_title('Dog')
fig.suptitle('Cat & Dog',size=30)
fig.tight_layout()
```



```
[12]: '''
      Through this random image selection and obseravtion,
      it is evident that shapes are not homogeneous;
      '''
```

```
[12]: '\nThrough this random image selection and obseravtion, \nit is evident that
      shapes are not homogeneous;\n'
```

0.0.4 Generators

```
[13]: from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

```
[14]: help(ImageDataGenerator)
```

Help on class ImageDataGenerator in module
tensorflow.python.keras.preprocessing.image:

```
class ImageDataGenerator(keras_preprocessing.image.image_data_generator.ImageDat
aGenerator)
```

```
| Generate batches of tensor image data with real-time data augmentation.
```

```
|
```

```
| The data will be looped over (in batches).
```

```
|
```

```

| Arguments:
|     featurewise_center: Boolean.
|         Set input mean to 0 over the dataset, feature-wise.
|     samplewise_center: Boolean. Set each sample mean to 0.
|     featurewise_std_normalization: Boolean.
|         Divide inputs by std of the dataset, feature-wise.
|     samplewise_std_normalization: Boolean. Divide each input by its std.
|     zca_epsilon: epsilon for ZCA whitening. Default is 1e-6.
|     zca_whitening: Boolean. Apply ZCA whitening.
|     rotation_range: Int. Degree range for random rotations.
|     width_shift_range: Float, 1-D array-like or int
|         - float: fraction of total width, if < 1, or pixels if >= 1.
|         - 1-D array-like: random elements from the array.
|         - int: integer number of pixels from interval
|             `(-width_shift_range, +width_shift_range)`
|         - With `width_shift_range=2` possible values
|             are integers `[-1, 0, +1]`,
|             same as with `width_shift_range=[-1, 0, +1]`,
|             while with `width_shift_range=1.0` possible values are floats
|             in the interval [-1.0, +1.0).
|     height_shift_range: Float, 1-D array-like or int
|         - float: fraction of total height, if < 1, or pixels if >= 1.
|         - 1-D array-like: random elements from the array.
|         - int: integer number of pixels from interval
|             `(-height_shift_range, +height_shift_range)`
|         - With `height_shift_range=2` possible values
|             are integers `[-1, 0, +1]`,
|             same as with `height_shift_range=[-1, 0, +1]`,
|             while with `height_shift_range=1.0` possible values are floats
|             in the interval [-1.0, +1.0).
|     brightness_range: Tuple or list of two floats. Range for picking
|         a brightness shift value from.
|     shear_range: Float. Shear Intensity
|         (Shear angle in counter-clockwise direction in degrees)
|     zoom_range: Float or [lower, upper]. Range for random zoom.
|         If a float, `[lower, upper] = [1-zoom_range, 1+zoom_range]`.
|     channel_shift_range: Float. Range for random channel shifts.
|     fill_mode: One of {"constant", "nearest", "reflect" or "wrap"}.
|         Default is 'nearest'.
|         Points outside the boundaries of the input are filled
|         according to the given mode:
|         - 'constant': kkkkkkkk|abcd|kkkkkkkk (cval=k)
|         - 'nearest':  aaaaaaaa|abcd|dddddddd
|         - 'reflect':  abcd dcba|abcd|dcba abcd
|         - 'wrap':    abcdabcd|abcd|abcdabcd
|     cval: Float or Int.
|         Value used for points outside the boundaries
|         when `fill_mode = "constant"`.

```

```

| horizontal_flip: Boolean. Randomly flip inputs horizontally.
| vertical_flip: Boolean. Randomly flip inputs vertically.
| rescale: rescaling factor. Defaults to None.
|     If None or 0, no rescaling is applied,
|     otherwise we multiply the data by the value provided
|     (after applying all other transformations).
| preprocessing_function: function that will be applied on each input.
|     The function will run after the image is resized and augmented.
|     The function should take one argument:
|     one image (Numpy tensor with rank 3),
|     and should output a Numpy tensor with the same shape.
| data_format: Image data format,
|     either "channels_first" or "channels_last".
|     "channels_last" mode means that the images should have shape
|     `(samples, height, width, channels)`,
|     "channels_first" mode means that the images should have shape
|     `(samples, channels, height, width)`.
|     It defaults to the `image_data_format` value found in your
|     Keras config file at `~/.keras/keras.json`.
|     If you never set it, then it will be "channels_last".
| validation_split: Float. Fraction of images reserved for validation
|     (strictly between 0 and 1).
| dtype: Dtype to use for the generated arrays.

```

Examples:

Example of using `.flow(x, y)`:

```

| ```python
| (x_train, y_train), (x_test, y_test) = cifar10.load_data()
| y_train = np_utils.to_categorical(y_train, num_classes)
| y_test = np_utils.to_categorical(y_test, num_classes)
| datagen = ImageDataGenerator(
|     featurewise_center=True,
|     featurewise_std_normalization=True,
|     rotation_range=20,
|     width_shift_range=0.2,
|     height_shift_range=0.2,
|     horizontal_flip=True)
| # compute quantities required for featurewise normalization
| # (std, mean, and principal components if ZCA whitening is applied)
| datagen.fit(x_train)
| # fits the model on batches with real-time data augmentation:
| model.fit(datagen.flow(x_train, y_train, batch_size=32),
|           steps_per_epoch=len(x_train) / 32, epochs=epochs)
| # here's a more "manual" example
| for e in range(epochs):
|     print('Epoch', e)

```

```

|     batches = 0
|     for x_batch, y_batch in datagen.flow(x_train, y_train, batch_size=32):
|         model.fit(x_batch, y_batch)
|         batches += 1
|         if batches >= len(x_train) / 32:
|             # we need to break the loop by hand because
|             # the generator loops indefinitely
|             break
|     ...

```

Example of using `.flow_from_directory(directory)`:`

```

|     ```python
|     train_datagen = ImageDataGenerator(
|         rescale=1./255,
|         shear_range=0.2,
|         zoom_range=0.2,
|         horizontal_flip=True)
|     test_datagen = ImageDataGenerator(rescale=1./255)
|     train_generator = train_datagen.flow_from_directory(
|         'data/train',
|         target_size=(150, 150),
|         batch_size=32,
|         class_mode='binary')
|     validation_generator = test_datagen.flow_from_directory(
|         'data/validation',
|         target_size=(150, 150),
|         batch_size=32,
|         class_mode='binary')
|     model.fit(
|         train_generator,
|         steps_per_epoch=2000,
|         epochs=50,
|         validation_data=validation_generator,
|         validation_steps=800)
|     ...

```

Example of transforming images and masks together.

```

|     ```python
|     # we create two instances with the same arguments
|     data_gen_args = dict(featurewise_center=True,
|         featurewise_std_normalization=True,
|         rotation_range=90,
|         width_shift_range=0.1,
|         height_shift_range=0.1,
|         zoom_range=0.2)
|     image_datagen = ImageDataGenerator(**data_gen_args)

```

```

| mask_datagen = ImageDataGenerator(**data_gen_args)
| # Provide the same seed and keyword arguments to the fit and flow methods
| seed = 1
| image_datagen.fit(images, augment=True, seed=seed)
| mask_datagen.fit(masks, augment=True, seed=seed)
| image_generator = image_datagen.flow_from_directory(
|     'data/images',
|     class_mode=None,
|     seed=seed)
| mask_generator = mask_datagen.flow_from_directory(
|     'data/masks',
|     class_mode=None,
|     seed=seed)
| # combine generators into one which yields image and masks
| train_generator = zip(image_generator, mask_generator)
| model.fit(
|     train_generator,
|     steps_per_epoch=2000,
|     epochs=50)
| ...
|
| Method resolution order:
|     ImageDataGenerator
|     keras_preprocessing.image.image_data_generator.ImageDataGenerator
|     builtins.object
|
| Methods defined here:
|
|     __init__(self, featurewise_center=False, samplewise_center=False,
featurewise_std_normalization=False, samplewise_std_normalization=False,
zca_whitening=False, zca_epsilon=1e-06, rotation_range=0, width_shift_range=0.0,
height_shift_range=0.0, brightness_range=None, shear_range=0.0, zoom_range=0.0,
channel_shift_range=0.0, fill_mode='nearest', cval=0.0, horizontal_flip=False,
vertical_flip=False, rescale=None, preprocessing_function=None,
data_format=None, validation_split=0.0, dtype=None)
|         Initialize self. See help(type(self)) for accurate signature.
|
|     flow(self, x, y=None, batch_size=32, shuffle=True, sample_weight=None,
seed=None, save_to_dir=None, save_prefix='', save_format='png', subset=None)
|         Takes data & label arrays, generates batches of augmented data.
|
|     Arguments:
|         x: Input data. Numpy array of rank 4 or a tuple. If tuple, the first
|             element should contain the images and the second element another
numpy
|             array or a list of numpy arrays that gets passed to the output
without
|             any modifications. Can be used to feed the model miscellaneous

```

```

data
|           along with the images. In case of grayscale data, the channels
axis of
|           the image array should have value 1, in case of RGB data, it
should
|           have value 3, and in case of RGBA data, it should have value 4.
|
|           y: Labels.
|           batch_size: Int (default: 32).
|           shuffle: Boolean (default: True).
|           sample_weight: Sample weights.
|           seed: Int (default: None).
|           save_to_dir: None or str (default: None). This allows you to
optionally
|           specify a directory to which to save the augmented pictures being
|           generated (useful for visualizing what you are doing).
|           save_prefix: Str (default: ''). Prefix to use for filenames of
saved
|           pictures (only relevant if `save_to_dir` is set).
|           save_format: one of "png", "jpeg"
|           (only relevant if `save_to_dir` is set). Default: "png".
|           subset: Subset of data ("training" or "validation") if
|           `validation_split` is set in `ImageDataGenerator`.
|
|
| Returns:
|
|     An `Iterator` yielding tuples of `(x, y)`
|     where `x` is a numpy array of image data
|     (in the case of a single image input) or a list
|     of numpy arrays (in the case with
|     additional inputs) and `y` is a numpy array
|     of corresponding labels. If 'sample_weight' is not None,
|     the yielded tuples are of the form `(x, y, sample_weight)`.
|     If `y` is None, only the numpy array `x` is returned.
|
|
|     flow_from_dataframe(self, dataframe, directory=None, x_col='filename',
y_col='class', weight_col=None, target_size=(256, 256), color_mode='rgb',
classes=None, class_mode='categorical', batch_size=32, shuffle=True, seed=None,
save_to_dir=None, save_prefix='', save_format='png', subset=None,
interpolation='nearest', validate_filenames=True, **kwargs)
|
|     Takes the dataframe and the path to a directory + generates batches.
|
|
|     The generated batches contain augmented/normalized data.
|
|
|     **A simple tutorial can be found **[here](
|                                     http://bit.ly/keras_flow_from_dataframe).
|
|
| Arguments:
|
|     dataframe: Pandas dataframe containing the filepaths relative to
|     `directory` (or absolute paths if `directory` is None) of the

```



```

images
|           in a string column. It should include other column/s
|           depending on the `class_mode`: - if `class_mode` is
`"categorical"`
|           (default value) it must include the `y_col` column with the
|           class/es of each image. Values in column can be
string/list/tuple
|           if a single class or list/tuple if multiple classes. - if
|           `class_mode` is `"binary"` or `"sparse"` it must include the
given
|           `y_col` column with class values as strings. - if `class_mode`
is
|           `"raw"` or `"multi_output"` it should contain the columns
|           specified in `y_col`. - if `class_mode` is `"input"` or `None`
no
|           extra column is needed.
|           directory: string, path to the directory to read images from. If
`None`,
|           data in `x_col` column should be absolute paths.
|           x_col: string, column in `dataframe` that contains the filenames (or
|           absolute paths if `directory` is `None`).
|           y_col: string or list, column/s in `dataframe` that has the target
data.
|           weight_col: string, column in `dataframe` that contains the sample
|           weights. Default: `None`.
|           target_size: tuple of integers `(height, width)`, default: `(256,
256)`.
|           The dimensions to which all images found will be resized.
|           color_mode: one of "grayscale", "rgb", "rgba". Default: "rgb".
Whether
|           the images will be converted to have 1 or 3 color channels.
|           classes: optional list of classes (e.g. `['dogs', 'cats']`). Default
is
|           None. If not provided, the list of classes will be automatically
|           inferred from the `y_col`, which will map to the label indices,
will
|           be alphanumeric). The dictionary containing the mapping from class
|           names to class indices can be obtained via the attribute
|           `class_indices`.
|           class_mode: one of "binary", "categorical", "input", "multi_output",
|           "raw", "sparse" or None. Default: "categorical".
|           Mode for yielding the targets:
|           - `"binary"`: 1D numpy array of binary labels,
|           - `"categorical"`: 2D numpy array of one-hot encoded labels.
|           Supports multi-label output.
|           - `"input"`: images identical to input images (mainly used to
work
|           with autoencoders),

```

```

|         - `"multi_output"`: list with the values of the different
columns,
|         - `"raw"`: numpy array of values in `y_col` column(s),
|         - `"sparse"`: 1D numpy array of integer labels, - `None`, no
targets
|         are returned (the generator will only yield batches of image
data,
|         which is useful to use in `model.predict()`).
|     batch_size: size of the batches of data (default: 32).
|     shuffle: whether to shuffle the data (default: True)
|     seed: optional random seed for shuffling and transformations.
|     save_to_dir: None or str (default: None). This allows you to
optionally
|         specify a directory to which to save the augmented pictures being
|         generated (useful for visualizing what you are doing).
|     save_prefix: str. Prefix to use for filenames of saved pictures
(only
|         relevant if `save_to_dir` is set).
|     save_format: one of "png", "jpeg"
|         (only relevant if `save_to_dir` is set). Default: "png".
|     subset: Subset of data (`"training"` or `"validation"`) if
|         `validation_split` is set in `ImageDataGenerator`.
|     interpolation: Interpolation method used to resample the image if
the
|         target size is different from that of the loaded image. Supported
|         methods are `"nearest"`, `"bilinear"`, and `"bicubic"`. If PIL
version
|         1.1.3 or newer is installed, `"lanczos"` is also supported. If PIL
|         version 3.4.0 or newer is installed, `"box"` and `"hamming"` are
also
|         supported. By default, `"nearest"` is used.
|     validate_filenames: Boolean, whether to validate image filenames in
|         `x_col`. If `True`, invalid images will be ignored. Disabling this
|         option can lead to speed-up in the execution of this function.
|         Defaults to `True`.
|     **kwargs: legacy arguments for raising deprecation warnings.
|
|     Returns:
|         A `DataFrameIterator` yielding tuples of `(x, y)`
|         where `x` is a numpy array containing a batch
|         of images with shape `(batch_size, *target_size, channels)`
|         and `y` is a numpy array of corresponding labels.
|
|     flow_from_directory(self, directory, target_size=(256, 256),
color_mode='rgb', classes=None, class_mode='categorical', batch_size=32,
shuffle=True, seed=None, save_to_dir=None, save_prefix='', save_format='png',
follow_links=False, subset=None, interpolation='nearest')
|         Takes the path to a directory & generates batches of augmented data.

```

|
 | Arguments:
 | directory: string, path to the target directory. It should contain
 one
 | subdirectory per class. Any PNG, JPG, BMP, PPM or TIF images
 inside
 | each of the subdirectories directory tree will be included in the
 | generator. See [this script](
 |
<https://gist.github.com/fchollet/0830affa1f7f19fd47b06d4cf89ed44d>)
 | for more details.
 | target_size: Tuple of integers `(height, width)`, defaults to `(256,
 | 256)`. The dimensions to which all images found will be resized.
 | color_mode: One of "grayscale", "rgb", "rgba". Default: "rgb".
 Whether
 | the images will be converted to have 1, 3, or 4 channels.
 | classes: Optional list of class subdirectories
 | (e.g. `['dogs', 'cats']`). Default: None. If not provided, the
 list
 | of classes will be automatically inferred from the
 subdirectory
 | names/structure under `directory`, where each subdirectory
 will be
 | treated as a different class (and the order of the classes,
 which
 | will map to the label indices, will be alphanumeric). The
 | dictionary containing the mapping from class names to class
 | indices can be obtained via the attribute `class_indices`.
 | class_mode: One of "categorical", "binary", "sparse",
 | "input", or None. Default: "categorical".
 | Determines the type of label arrays that are returned: -
 | "categorical" will be 2D one-hot encoded labels, - "binary"
 will
 | be 1D binary labels, "sparse" will be 1D integer labels, -
 "input"
 | will be images identical to input images (mainly used to work
 with
 | autoencoders). - If None, no labels are returned (the
 generator
 | will only yield batches of image data, which is useful to use
 with
 | `model.predict()`). Please note that in case of
 | class_mode None, the data still needs to reside in a
 subdirectory
 | of `directory` for it to work correctly.
 | batch_size: Size of the batches of data (default: 32).
 | shuffle: Whether to shuffle the data (default: True) If set to
 False,

```

|         sorts the data in alphanumeric order.
|         seed: Optional random seed for shuffling and transformations.
|         save_to_dir: None or str (default: None). This allows you to
optionally
|         specify a directory to which to save the augmented pictures being
|         generated (useful for visualizing what you are doing).
|         save_prefix: Str. Prefix to use for filenames of saved pictures
(only
|         relevant if `save_to_dir` is set).
|         save_format: One of "png", "jpeg"
|         (only relevant if `save_to_dir` is set). Default: "png".
|         follow_links: Whether to follow symlinks inside
|         class subdirectories (default: False).
|         subset: Subset of data (`"training"` or `"validation"`) if
|         `validation_split` is set in `ImageDataGenerator`.
|         interpolation: Interpolation method used to resample the image if
the
|         target size is different from that of the loaded image. Supported
|         methods are `"nearest"`, `"bilinear"`, and `"bicubic"`. If PIL
version
|         1.1.3 or newer is installed, `"lanczos"` is also supported. If PIL
|         version 3.4.0 or newer is installed, `"box"` and `"hamming"` are
also
|         supported. By default, `"nearest"` is used.
|
|     Returns:
|         A `DirectoryIterator` yielding tuples of `(x, y)`
|         where `x` is a numpy array containing a batch
|         of images with shape `(batch_size, *target_size, channels)`
|         and `y` is a numpy array of corresponding labels.
|
|     -----
|     Methods inherited from
keras_preprocessing.image.image_data_generator.ImageDataGenerator:
|
|     apply_transform(self, x, transform_parameters)
|         Applies a transformation to an image according to given parameters.
|
|     # Arguments
|         x: 3D tensor, single image.
|         transform_parameters: Dictionary with string - parameter pairs
|         describing the transformation.
|         Currently, the following parameters
|         from the dictionary are used:
|         - `theta`: Float. Rotation angle in degrees.
|         - `tx`: Float. Shift in the x direction.
|         - `ty`: Float. Shift in the y direction.
|         - `shear`: Float. Shear angle in degrees.

```

```

|         - ``'zx'``: Float. Zoom in the x direction.
|         - ``'zy'``: Float. Zoom in the y direction.
|         - ``'flip_horizontal'``: Boolean. Horizontal flip.
|         - ``'flip_vertical'``: Boolean. Vertical flip.
|         - ``'channel_shift_intensity'``: Float. Channel shift intensity.
|         - ``'brightness'``: Float. Brightness shift intensity.
|
|     # Returns
|         A transformed version of the input (same shape).
|
| fit(self, x, augment=False, rounds=1, seed=None)
|     Fits the data generator to some sample data.
|
|     This computes the internal data stats related to the
|     data-dependent transformations, based on an array of sample data.
|
|     Only required if ``featurewise_center`` or
|     ``featurewise_std_normalization`` or ``zca_whitening`` are set to True.
|
|     When ``rescale`` is set to a value, rescaling is applied to
|     sample data before computing the internal data stats.
|
|     # Arguments
|         x: Sample data. Should have rank 4.
|             In case of grayscale data,
|             the channels axis should have value 1, in case
|             of RGB data, it should have value 3, and in case
|             of RGBA data, it should have value 4.
|         augment: Boolean (default: False).
|             Whether to fit on randomly augmented samples.
|         rounds: Int (default: 1).
|             If using data augmentation (``augment=True``),
|             this is how many augmentation passes over the data to use.
|         seed: Int (default: None). Random seed.
|
| get_random_transform(self, img_shape, seed=None)
|     Generates random parameters for a transformation.
|
|     # Arguments
|         seed: Random seed.
|         img_shape: Tuple of integers.
|             Shape of the image that is transformed.
|
|     # Returns
|         A dictionary containing randomly chosen parameters describing the
|         transformation.
|
| random_transform(self, x, seed=None)

```

```

|     Applies a random transformation to an image.
|
|     # Arguments
|         x: 3D tensor, single image.
|         seed: Random seed.
|
|     # Returns
|         A randomly transformed version of the input (same shape).
|
| standardize(self, x)
|     Applies the normalization configuration in-place to a batch of inputs.
|
|     `x` is changed in-place since the function is mainly used internally
|     to standardize images and feed them to your network. If a copy of `x`
|     would be created instead it would have a significant performance cost.
|     If you want to apply this method without changing the input in-place
|     you can call the method creating a copy before:
|
|     standardize(np.copy(x))
|
|     # Arguments
|         x: Batch of inputs to be normalized.
|
|     # Returns
|         The inputs, normalized.
|
| -----
| Data descriptors inherited from
keras_preprocessing.image.image_data_generator.ImageDataGenerator:
|
|     __dict__
|         dictionary for instance variables (if defined)
|
|     __weakref__
|         list of weak references to the object (if defined)

```

```

[15]: train_gen = ImageDataGenerator(rescale=1/255.)
      val_gen = ImageDataGenerator(rescale=1/255.)

```

```

[16]: # both cats and dogs combined
      train_generator = train_gen.
      ↳flow_from_directory(train_dir,target_size=(128,128),batch_size=64,shuffle=True,class_mode='
      validation_generator = val_gen.
      ↳flow_from_directory(valid_dir,target_size=(128,128),batch_size=64,shuffle=False,class_mode=

```

Found 18743 images belonging to 2 classes.
Found 6251 images belonging to 2 classes.

Generators does almost everything that we need in a robust way by just parameterizing what we need. in this case because the image shapes where Heterogeneous in nature, we took a target size of 128,128 to transform every image into that size.

```
[17]: total_data = train_generator.samples + validation_generator.samples
print('Classes Include : ',train_generator.class_indices)
print('Train Samples : ', train_generator.samples)
print('Validation/Test Samples : ', validation_generator.samples)
print('Total Samples present (cats+dogs) : ', total_data)
print('Transformed Image shape : ',train_generator.image_shape)
print('Color Channels : (RGB)', train_generator.image_shape[2])
```

```
Classes Include : {'CAT': 0, 'DOG': 1}
Train Samples : 18743
Validation/Test Samples : 6251
Total Samples present (cats+dogs) : 24994
Transformed Image shape : (128, 128, 3)
Color Channels : (RGB) 3
```

0.0.5 Ordinary Convolutional Neural Network Method

```
[18]: import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Conv2D, MaxPooling2D, Flatten,
↳BatchNormalization, Activation, Dropout
```

```
[34]: input_shape = train_generator.image_shape

net = Sequential()

#convnet

net.add(Conv2D(32, kernel_size=3,input_shape = train_generator.image_shape,
↳activation='relu'))
net.add(MaxPooling2D(2,2))

net.add(Conv2D(64,(3,3), activation='relu'))
net.add(MaxPooling2D(2,2))

net.add(Conv2D(128, (3,3), activation='relu'))
net.add(MaxPooling2D(2,2))

net.add(Flatten())

net.add(Dense(512, activation='relu'))
net.add(Dense(1,activation='sigmoid'))
```

```
[35]: net.summary()
```

```
Model: "sequential_1"
```

Layer (type)	Output Shape	Param #
conv2d_3 (Conv2D)	(None, 126, 126, 32)	896
max_pooling2d_3 (MaxPooling2D)	(None, 63, 63, 32)	0
conv2d_4 (Conv2D)	(None, 61, 61, 64)	18496
max_pooling2d_4 (MaxPooling2D)	(None, 30, 30, 64)	0
conv2d_5 (Conv2D)	(None, 28, 28, 128)	73856
max_pooling2d_5 (MaxPooling2D)	(None, 14, 14, 128)	0
flatten_1 (Flatten)	(None, 25088)	0
dense_2 (Dense)	(None, 512)	12845568
dense_3 (Dense)	(None, 1)	513

Total params: 12,939,329
Trainable params: 12,939,329
Non-trainable params: 0

```
[36]: from tensorflow.keras.callbacks import EarlyStopping  
adam = tf.keras.optimizers.Adam(lr=0.001)
```

```
[37]: net.compile(optimizer=adam, loss='binary_crossentropy', metrics=['accuracy'])
```

```
[38]: early_stop = EarlyStopping(patience=3)
```

```
[39]: import time  
  
start = time.perf_counter()  
  
perf = net.fit_generator(train_generator, epochs=15,  
    ↳ callbacks=[early_stop], validation_data=validation_generator)  
  
elapsed = time.perf_counter() - start  
  
print('Elapsed {}'.format(elapsed/60))
```



```

/usr/local/lib/python3.6/dist-packages/tensorflow/python/keras/engine/training.py:1844: UserWarning:
`Model.fit_generator` is deprecated and will be removed in a future version.
Please use `Model.fit`, which supports generators.
  warnings.warn("`Model.fit_generator` is deprecated and '
Epoch 1/15
 2/293 [...] - ETA: 1:07 - loss: 1.2549 - accuracy:
0.4570

/usr/local/lib/python3.6/dist-packages/PIL/TiffImagePlugin.py:770: UserWarning:
Possibly corrupt EXIF data. Expecting to read 32 bytes but only got 0. Skipping
tag 270
  " Skipping tag %s" % (size, len(data), tag)
/usr/local/lib/python3.6/dist-packages/PIL/TiffImagePlugin.py:770: UserWarning:
Possibly corrupt EXIF data. Expecting to read 5 bytes but only got 0. Skipping
tag 271
  " Skipping tag %s" % (size, len(data), tag)
/usr/local/lib/python3.6/dist-packages/PIL/TiffImagePlugin.py:770: UserWarning:
Possibly corrupt EXIF data. Expecting to read 8 bytes but only got 0. Skipping
tag 272
  " Skipping tag %s" % (size, len(data), tag)
/usr/local/lib/python3.6/dist-packages/PIL/TiffImagePlugin.py:770: UserWarning:
Possibly corrupt EXIF data. Expecting to read 8 bytes but only got 0. Skipping
tag 282
  " Skipping tag %s" % (size, len(data), tag)
/usr/local/lib/python3.6/dist-packages/PIL/TiffImagePlugin.py:770: UserWarning:
Possibly corrupt EXIF data. Expecting to read 8 bytes but only got 0. Skipping
tag 283
  " Skipping tag %s" % (size, len(data), tag)
/usr/local/lib/python3.6/dist-packages/PIL/TiffImagePlugin.py:770: UserWarning:
Possibly corrupt EXIF data. Expecting to read 20 bytes but only got 0. Skipping
tag 306
  " Skipping tag %s" % (size, len(data), tag)
/usr/local/lib/python3.6/dist-packages/PIL/TiffImagePlugin.py:770: UserWarning:
Possibly corrupt EXIF data. Expecting to read 48 bytes but only got 0. Skipping
tag 532
  " Skipping tag %s" % (size, len(data), tag)
/usr/local/lib/python3.6/dist-packages/PIL/TiffImagePlugin.py:788: UserWarning:
Corrupt EXIF data. Expecting to read 2 bytes but only got 0.
  warnings.warn(str(msg))

293/293 [=====] - 92s 312ms/step - loss: 0.6959 -
accuracy: 0.5951 - val_loss: 0.5447 - val_accuracy: 0.7207
Epoch 2/15
293/293 [=====] - 91s 310ms/step - loss: 0.5172 -
accuracy: 0.7472 - val_loss: 0.4482 - val_accuracy: 0.7906
Epoch 3/15
293/293 [=====] - 92s 313ms/step - loss: 0.4282 -

```

```

accuracy: 0.8010 - val_loss: 0.4345 - val_accuracy: 0.7903
Epoch 4/15
293/293 [=====] - 91s 310ms/step - loss: 0.3586 -
accuracy: 0.8409 - val_loss: 0.3913 - val_accuracy: 0.8219
Epoch 5/15
293/293 [=====] - 92s 313ms/step - loss: 0.3031 -
accuracy: 0.8719 - val_loss: 0.3582 - val_accuracy: 0.8416
Epoch 6/15
293/293 [=====] - 92s 314ms/step - loss: 0.2391 -
accuracy: 0.9005 - val_loss: 0.4140 - val_accuracy: 0.8322
Epoch 7/15
293/293 [=====] - 91s 310ms/step - loss: 0.1694 -
accuracy: 0.9329 - val_loss: 0.3971 - val_accuracy: 0.8440
Epoch 8/15
293/293 [=====] - 91s 311ms/step - loss: 0.1156 -
accuracy: 0.9560 - val_loss: 0.4748 - val_accuracy: 0.8375
Elapsed 12.185325518533329

```

```
[40]: loss, acc = net.evaluate_generator(validation_generator)
```

```

/usr/local/lib/python3.6/dist-
packages/tensorflow/python/keras/engine/training.py:1877: UserWarning:
`Model.evaluate_generator` is deprecated and will be removed in a future
version. Please use `Model.evaluate`, which supports generators.
  warnings.warn("`Model.evaluate_generator` is deprecated and ")

```

```

[44]: # validation Accuracy
print('Validation Accuracy - {:.2f}%'.format(acc*100))
print('validation loss - {:.2f}'.format(loss))

```

```

Validation Accuracy - 83.75%
validation loss - 0.47

```

```
[43]: history = pd.DataFrame(perf.history)
```


```
[43]:
```

```

[45]: plt.style.use('seaborn-whitegrid')
plt.plot(history['val_loss'], ls='-', color='blue', label='val_loss')
plt.plot(history['loss'], ls='--', color='orange', label='loss')
plt.legend()

```

```
[45]: <matplotlib.legend.Legend at 0x7f234a9e0198>
```



```
[50]: plt.style.use('seaborn-whitegrid')
plt.plot(history['val_accuracy'], ls='-',color='blue',label='val_loss')
plt.plot(history['accuracy'],ls='--',color='orange',label='loss' )
plt.legend(loc='upper left')
```

[50]: <matplotlib.legend.Legend at 0x7f234a6bb4a8>



```
[ ]: # Classwise Evaluation
```

```
[52]: pred_probs = net.predict_generator(validation_generator)
      preds = pred_probs > 0.5
```

```
/usr/local/lib/python3.6/dist-packages/tensorflow/python/keras/engine/training.py:1905: UserWarning:
`Model.predict_generator` is deprecated and will be removed in a future version.
Please use `Model.predict`, which supports generators.
  warnings.warn("`Model.predict_generator` is deprecated and "
```

```
[53]: preds
```

```
[53]: array([[ True],
           [False],
           [False],
           ...,
           [ True],
           [ True],
           [ True]])
```

```
[54]: from sklearn.metrics import classification_report, confusion_matrix
```

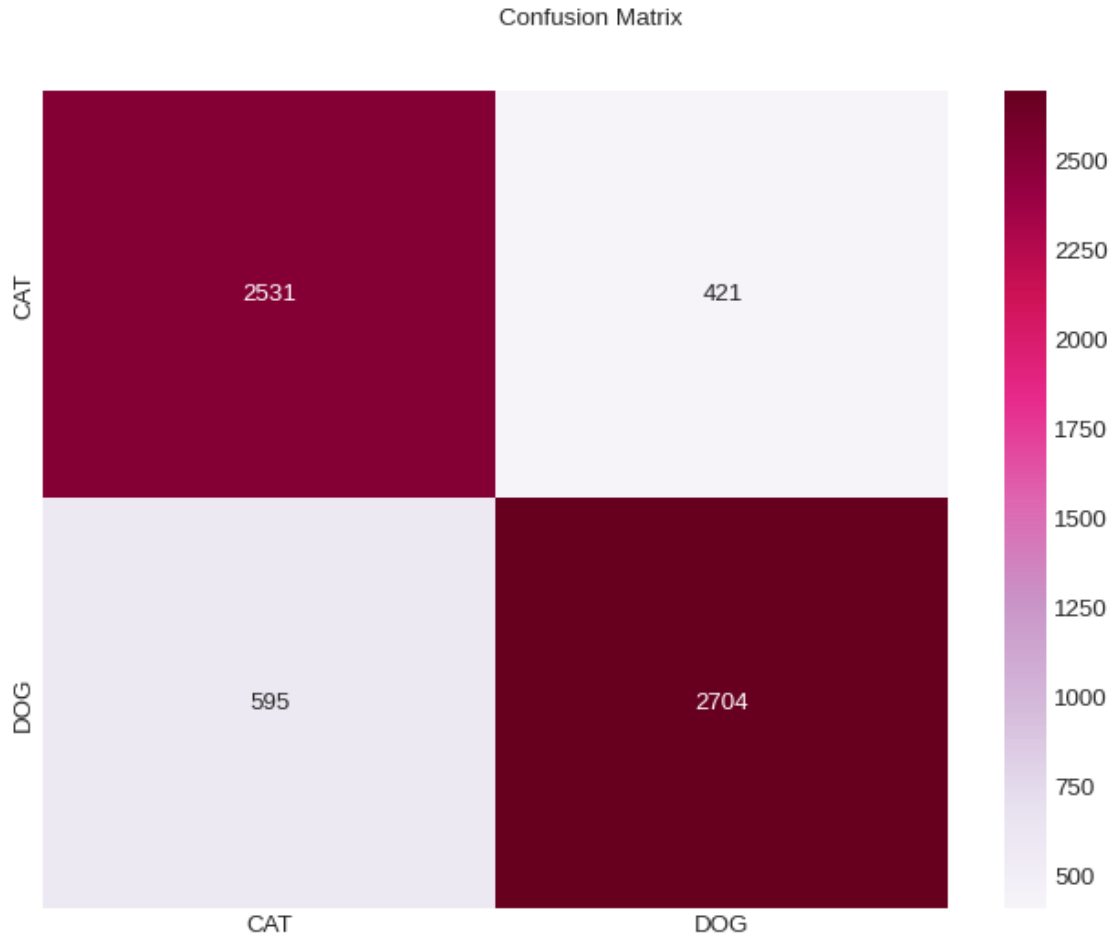
```
[55]: print(classification_report(preds, validation_generator.classes))
```

	precision	recall	f1-score	support
False	0.81	0.86	0.83	2952
True	0.87	0.82	0.84	3299
accuracy			0.84	6251
macro avg	0.84	0.84	0.84	6251
weighted avg	0.84	0.84	0.84	6251

```
[ ]:
```

```
[56]: plt.figure(figsize=(10,7))
plt.suptitle('Confusion Matrix')
confmat = pd.DataFrame(confusion_matrix(preds, validation_generator.classes),
↳ columns=validation_generator.class_indices.keys())
confmat.index = validation_generator.class_indices.keys()
sns.heatmap(confmat, annot=True, fmt='d', cmap='PuRd')
```

```
[56]: <matplotlib.axes._subplots.AxesSubplot at 0x7f23ac2deba8>
```



The Classic case of Misclassifying dogs as cats and cats as dogs. the latter one is mostly happening i.e. cats as dogs classification

```
[94]: net.save('CovnNet(cats&dogs).h5')
```

0.0.6 Inference

While the training Set Accuracy steadily increases till it reaches 90 the validation accuracy is stall at 80's, it is clearly the case of overfitting or high variance scenario.

```
[83]: from tensorflow.keras.applications import MobileNetV2
```

```
[84]: train_generator.image_shape
```

```
[84]: (128, 128, 3)
```

```
[85]: conv_base = MobileNetV2(input_shape=train_generator.image_shape,
    ↳ include_top=False, weights='imagenet')
```

```
[86]: conv_base.trainable
```

```
[86]: True
```

```
[87]: # freezing layers  
conv_base.trainable= False
```

```
[88]: conv_base.summary()
```

```
Model: "mobilenetv2_1.00_128"
```

```
-----  
-----  
Layer (type)                Output Shape          Param #   Connected to  
=====
```

input_2 (InputLayer)	[(None, 128, 128, 3)	0	

Conv1 (Conv2D)	(None, 64, 64, 32)	864	input_2[0][0]

bn_Conv1 (BatchNormalization)	(None, 64, 64, 32)	128	Conv1[0][0]

Conv1_relu (ReLU)	(None, 64, 64, 32)	0	bn_Conv1[0][0]

expanded_conv_depthwise (Depthw	(None, 64, 64, 32)	288	
Conv1_relu[0][0]			

expanded_conv_depthwise_BN (Bat	(None, 64, 64, 32)	128	
expanded_conv_depthwise[0][0]			

expanded_conv_depthwise_relu (R	(None, 64, 64, 32)	0	
expanded_conv_depthwise_BN[0][0]			

expanded_conv_project (Conv2D)	(None, 64, 64, 16)	512	
expanded_conv_depthwise_relu[0][0]			

expanded_conv_project_BN (Batch	(None, 64, 64, 16)	64	
expanded_conv_project[0][0]			
----- -----			

```

block_1_expand (Conv2D)          (None, 64, 64, 96)    1536
expanded_conv_project_BN[0][0]

```

```

-----
block_1_expand_BN (BatchNormali (None, 64, 64, 96)    384
block_1_expand[0][0]

```

```

-----
block_1_expand_relu (ReLU)       (None, 64, 64, 96)    0
block_1_expand_BN[0][0]

```

```

-----
block_1_pad (ZeroPadding2D)      (None, 65, 65, 96)    0
block_1_expand_relu[0][0]

```

```

-----
block_1_depthwise (DepthwiseCon (None, 32, 32, 96)    864
block_1_pad[0][0]

```

```

-----
block_1_depthwise_BN (BatchNorm (None, 32, 32, 96)    384
block_1_depthwise[0][0]

```

```

-----
block_1_depthwise_relu (ReLU)    (None, 32, 32, 96)    0
block_1_depthwise_BN[0][0]

```

```

-----
block_1_project (Conv2D)         (None, 32, 32, 24)    2304
block_1_depthwise_relu[0][0]

```

```

-----
block_1_project_BN (BatchNormal (None, 32, 32, 24)    96
block_1_project[0][0]

```

```

-----
block_2_expand (Conv2D)          (None, 32, 32, 144)   3456
block_1_project_BN[0][0]

```

```

-----
block_2_expand_BN (BatchNormali (None, 32, 32, 144)   576
block_2_expand[0][0]

```

```

-----
block_2_expand_relu (ReLU)       (None, 32, 32, 144)   0
block_2_expand_BN[0][0]

```


block_2_depthwise (DepthwiseCon (None, 32, 32, 144) 1296
block_2_expand_relu[0][0]

block_2_depthwise_BN (BatchNorm (None, 32, 32, 144) 576
block_2_depthwise[0][0]

block_2_depthwise_relu (ReLU) (None, 32, 32, 144) 0
block_2_depthwise_BN[0][0]

block_2_project (Conv2D) (None, 32, 32, 24) 3456
block_2_depthwise_relu[0][0]

block_2_project_BN (BatchNormal (None, 32, 32, 24) 96
block_2_project[0][0]

block_2_add (Add) (None, 32, 32, 24) 0
block_1_project_BN[0][0]
block_2_project_BN[0][0]

block_3_expand (Conv2D) (None, 32, 32, 144) 3456
block_2_add[0][0]

block_3_expand_BN (BatchNormali (None, 32, 32, 144) 576
block_3_expand[0][0]

block_3_expand_relu (ReLU) (None, 32, 32, 144) 0
block_3_expand_BN[0][0]

block_3_pad (ZeroPadding2D) (None, 33, 33, 144) 0
block_3_expand_relu[0][0]

block_3_depthwise (DepthwiseCon (None, 16, 16, 144) 1296
block_3_pad[0][0]

block_3_depthwise_BN (BatchNorm (None, 16, 16, 144) 576
block_3_depthwise[0][0]

```

-----
block_3_depthwise_relu (ReLU)      (None, 16, 16, 144)  0
block_3_depthwise_BN[0][0]
-----

-----
block_3_project (Conv2D)            (None, 16, 16, 32)   4608
block_3_depthwise_relu[0][0]
-----

-----
block_3_project_BN (BatchNormal (None, 16, 16, 32)   128
block_3_project[0][0]
-----

-----
block_4_expand (Conv2D)             (None, 16, 16, 192)  6144
block_3_project_BN[0][0]
-----

-----
block_4_expand_BN (BatchNormali (None, 16, 16, 192)  768
block_4_expand[0][0]
-----

-----
block_4_expand_relu (ReLU)          (None, 16, 16, 192)  0
block_4_expand_BN[0][0]
-----

-----
block_4_depthwise (DepthwiseCon (None, 16, 16, 192)  1728
block_4_expand_relu[0][0]
-----

-----
block_4_depthwise_BN (BatchNorm (None, 16, 16, 192)  768
block_4_depthwise[0][0]
-----

-----
block_4_depthwise_relu (ReLU)       (None, 16, 16, 192)  0
block_4_depthwise_BN[0][0]
-----

-----
block_4_project (Conv2D)            (None, 16, 16, 32)   6144
block_4_depthwise_relu[0][0]
-----

-----
block_4_project_BN (BatchNormal (None, 16, 16, 32)   128
block_4_project[0][0]
-----

-----
block_4_add (Add)                   (None, 16, 16, 32)   0
block_3_project_BN[0][0]
block_4_project_BN[0][0]

```

```

-----
-----
block_5_expand (Conv2D)          (None, 16, 16, 192)  6144
block_4_add[0][0]
-----
-----
block_5_expand_BN (BatchNormali (None, 16, 16, 192)  768
block_5_expand[0][0]
-----
-----
block_5_expand_relu (ReLU)       (None, 16, 16, 192)  0
block_5_expand_BN[0][0]
-----
-----
block_5_depthwise (DepthwiseCon (None, 16, 16, 192)  1728
block_5_expand_relu[0][0]
-----
-----
block_5_depthwise_BN (BatchNorm (None, 16, 16, 192)  768
block_5_depthwise[0][0]
-----
-----
block_5_depthwise_relu (ReLU)    (None, 16, 16, 192)  0
block_5_depthwise_BN[0][0]
-----
-----
block_5_project (Conv2D)         (None, 16, 16, 32)   6144
block_5_depthwise_relu[0][0]
-----
-----
block_5_project_BN (BatchNormal (None, 16, 16, 32)   128
block_5_project[0][0]
-----
-----
block_5_add (Add)                (None, 16, 16, 32)   0
block_4_add[0][0]
block_5_project_BN[0][0]
-----
-----
block_6_expand (Conv2D)          (None, 16, 16, 192)  6144
block_5_add[0][0]
-----
-----
block_6_expand_BN (BatchNormali (None, 16, 16, 192)  768
block_6_expand[0][0]
-----
-----
block_6_expand_relu (ReLU)       (None, 16, 16, 192)  0

```

block_6_expand_BN[0][0]

block_6_pad (ZeroPadding2D) (None, 17, 17, 192) 0
block_6_expand_relu[0][0]

block_6_depthwise (DepthwiseCon (None, 8, 8, 192) 1728
block_6_pad[0][0]

block_6_depthwise_BN (BatchNorm (None, 8, 8, 192) 768
block_6_depthwise[0][0]

block_6_depthwise_relu (ReLU) (None, 8, 8, 192) 0
block_6_depthwise_BN[0][0]

block_6_project (Conv2D) (None, 8, 8, 64) 12288
block_6_depthwise_relu[0][0]

block_6_project_BN (BatchNormal (None, 8, 8, 64) 256
block_6_project[0][0]

block_7_expand (Conv2D) (None, 8, 8, 384) 24576
block_6_project_BN[0][0]

block_7_expand_BN (BatchNormali (None, 8, 8, 384) 1536
block_7_expand[0][0]

block_7_expand_relu (ReLU) (None, 8, 8, 384) 0
block_7_expand_BN[0][0]

block_7_depthwise (DepthwiseCon (None, 8, 8, 384) 3456
block_7_expand_relu[0][0]

block_7_depthwise_BN (BatchNorm (None, 8, 8, 384) 1536
block_7_depthwise[0][0]

block_7_depthwise_relu (ReLU) (None, 8, 8, 384) 0

block_7_depthwise_BN[0][0]

block_7_project (Conv2D) (None, 8, 8, 64) 24576
block_7_depthwise_relu[0][0]

block_7_project_BN (BatchNormal (None, 8, 8, 64) 256
block_7_project[0][0]

block_7_add (Add) (None, 8, 8, 64) 0
block_6_project_BN[0][0]
block_7_project_BN[0][0]

block_8_expand (Conv2D) (None, 8, 8, 384) 24576
block_7_add[0][0]

block_8_expand_BN (BatchNormali (None, 8, 8, 384) 1536
block_8_expand[0][0]

block_8_expand_relu (ReLU) (None, 8, 8, 384) 0
block_8_expand_BN[0][0]

block_8_depthwise (DepthwiseCon (None, 8, 8, 384) 3456
block_8_expand_relu[0][0]

block_8_depthwise_BN (BatchNorm (None, 8, 8, 384) 1536
block_8_depthwise[0][0]

block_8_depthwise_relu (ReLU) (None, 8, 8, 384) 0
block_8_depthwise_BN[0][0]

block_8_project (Conv2D) (None, 8, 8, 64) 24576
block_8_depthwise_relu[0][0]

block_8_project_BN (BatchNormal (None, 8, 8, 64) 256
block_8_project[0][0]

block_8_add (Add)	(None, 8, 8, 64)	0
block_7_add[0][0]		
block_8_project_BN[0][0]		

block_9_expand (Conv2D)	(None, 8, 8, 384)	24576
block_8_add[0][0]		

block_9_expand_BN (BatchNormali	(None, 8, 8, 384)	1536
block_9_expand[0][0]		

block_9_expand_relu (ReLU)	(None, 8, 8, 384)	0
block_9_expand_BN[0][0]		

block_9_depthwise (DepthwiseCon	(None, 8, 8, 384)	3456
block_9_expand_relu[0][0]		

block_9_depthwise_BN (BatchNorm	(None, 8, 8, 384)	1536
block_9_depthwise[0][0]		

block_9_depthwise_relu (ReLU)	(None, 8, 8, 384)	0
block_9_depthwise_BN[0][0]		

block_9_project (Conv2D)	(None, 8, 8, 64)	24576
block_9_depthwise_relu[0][0]		

block_9_project_BN (BatchNormal	(None, 8, 8, 64)	256
block_9_project[0][0]		

block_9_add (Add)	(None, 8, 8, 64)	0
block_8_add[0][0]		
block_9_project_BN[0][0]		

block_10_expand (Conv2D)	(None, 8, 8, 384)	24576
block_9_add[0][0]		

block_10_expand_BN (BatchNormal	(None, 8, 8, 384)	1536
block_10_expand[0][0]		

```

-----
-----
block_10_expand_relu (ReLU)      (None, 8, 8, 384)      0
block_10_expand_BN[0][0]
-----
-----
block_10_depthwise (DepthwiseCo (None, 8, 8, 384)      3456
block_10_expand_relu[0][0]
-----
-----
block_10_depthwise_BN (BatchNor (None, 8, 8, 384)      1536
block_10_depthwise[0][0]
-----
-----
block_10_depthwise_relu (ReLU)   (None, 8, 8, 384)      0
block_10_depthwise_BN[0][0]
-----
-----
block_10_project (Conv2D)         (None, 8, 8, 96)       36864
block_10_depthwise_relu[0][0]
-----
-----
block_10_project_BN (BatchNorma (None, 8, 8, 96)       384
block_10_project[0][0]
-----
-----
block_11_expand (Conv2D)          (None, 8, 8, 576)      55296
block_10_project_BN[0][0]
-----
-----
block_11_expand_BN (BatchNormal (None, 8, 8, 576)      2304
block_11_expand[0][0]
-----
-----
block_11_expand_relu (ReLU)      (None, 8, 8, 576)      0
block_11_expand_BN[0][0]
-----
-----
block_11_depthwise (DepthwiseCo (None, 8, 8, 576)      5184
block_11_expand_relu[0][0]
-----
-----
block_11_depthwise_BN (BatchNor (None, 8, 8, 576)      2304
block_11_depthwise[0][0]
-----
-----
block_11_depthwise_relu (ReLU)   (None, 8, 8, 576)      0
block_11_depthwise_BN[0][0]

```

block_11_project (Conv2D)	(None, 8, 8, 96)	55296
block_11_depthwise_relu[0][0]		
block_11_project_BN (BatchNorma	(None, 8, 8, 96)	384
block_11_project[0][0]		
block_11_add (Add)	(None, 8, 8, 96)	0
block_10_project_BN[0][0]		
block_11_project_BN[0][0]		
block_12_expand (Conv2D)	(None, 8, 8, 576)	55296
block_11_add[0][0]		
block_12_expand_BN (BatchNormal	(None, 8, 8, 576)	2304
block_12_expand[0][0]		
block_12_expand_relu (ReLU)	(None, 8, 8, 576)	0
block_12_expand_BN[0][0]		
block_12_depthwise (DepthwiseCo	(None, 8, 8, 576)	5184
block_12_expand_relu[0][0]		
block_12_depthwise_BN (BatchNor	(None, 8, 8, 576)	2304
block_12_depthwise[0][0]		
block_12_depthwise_relu (ReLU)	(None, 8, 8, 576)	0
block_12_depthwise_BN[0][0]		
block_12_project (Conv2D)	(None, 8, 8, 96)	55296
block_12_depthwise_relu[0][0]		
block_12_project_BN (BatchNorma	(None, 8, 8, 96)	384
block_12_project[0][0]		
block_12_add (Add)	(None, 8, 8, 96)	0


```

block_11_add[0][0]
block_12_project_BN[0][0]
-----
-----
block_13_expand (Conv2D)          (None, 8, 8, 576)    55296
block_12_add[0][0]
-----
-----
block_13_expand_BN (BatchNormal (None, 8, 8, 576)    2304
block_13_expand[0][0]
-----
-----
block_13_expand_relu (ReLU)       (None, 8, 8, 576)    0
block_13_expand_BN[0][0]
-----
-----
block_13_pad (ZeroPadding2D)      (None, 9, 9, 576)    0
block_13_expand_relu[0][0]
-----
-----
block_13_depthwise (DepthwiseCo (None, 4, 4, 576)    5184
block_13_pad[0][0]
-----
-----
block_13_depthwise_BN (BatchNor (None, 4, 4, 576)    2304
block_13_depthwise[0][0]
-----
-----
block_13_depthwise_relu (ReLU)    (None, 4, 4, 576)    0
block_13_depthwise_BN[0][0]
-----
-----
block_13_project (Conv2D)          (None, 4, 4, 160)    92160
block_13_depthwise_relu[0][0]
-----
-----
block_13_project_BN (BatchNorma (None, 4, 4, 160)    640
block_13_project[0][0]
-----
-----
block_14_expand (Conv2D)          (None, 4, 4, 960)    153600
block_13_project_BN[0][0]
-----
-----
block_14_expand_BN (BatchNormal (None, 4, 4, 960)    3840
block_14_expand[0][0]
-----
-----

```

block_14_expand_relu (ReLU)	(None, 4, 4, 960)	0
block_14_expand_BN[0][0]		

block_14_depthwise (DepthwiseCo	(None, 4, 4, 960)	8640
block_14_expand_relu[0][0]		

block_14_depthwise_BN (BatchNor	(None, 4, 4, 960)	3840
block_14_depthwise[0][0]		

block_14_depthwise_relu (ReLU)	(None, 4, 4, 960)	0
block_14_depthwise_BN[0][0]		

block_14_project (Conv2D)	(None, 4, 4, 160)	153600
block_14_depthwise_relu[0][0]		

block_14_project_BN (BatchNorma	(None, 4, 4, 160)	640
block_14_project[0][0]		

block_14_add (Add)	(None, 4, 4, 160)	0
block_13_project_BN[0][0]		
block_14_project_BN[0][0]		

block_15_expand (Conv2D)	(None, 4, 4, 960)	153600
block_14_add[0][0]		

block_15_expand_BN (BatchNormal	(None, 4, 4, 960)	3840
block_15_expand[0][0]		

block_15_expand_relu (ReLU)	(None, 4, 4, 960)	0
block_15_expand_BN[0][0]		

block_15_depthwise (DepthwiseCo	(None, 4, 4, 960)	8640
block_15_expand_relu[0][0]		

block_15_depthwise_BN (BatchNor	(None, 4, 4, 960)	3840
block_15_depthwise[0][0]		

```

-----
block_15_depthwise_relu (ReLU) (None, 4, 4, 960) 0
block_15_depthwise_BN[0][0]
-----

-----
block_15_project (Conv2D) (None, 4, 4, 160) 153600
block_15_depthwise_relu[0][0]
-----

-----
block_15_project_BN (BatchNorma (None, 4, 4, 160) 640
block_15_project[0][0]
-----

-----
block_15_add (Add) (None, 4, 4, 160) 0
block_14_add[0][0]
block_15_project_BN[0][0]
-----

-----
block_16_expand (Conv2D) (None, 4, 4, 960) 153600
block_15_add[0][0]
-----

-----
block_16_expand_BN (BatchNormal (None, 4, 4, 960) 3840
block_16_expand[0][0]
-----

-----
block_16_expand_relu (ReLU) (None, 4, 4, 960) 0
block_16_expand_BN[0][0]
-----

-----
block_16_depthwise (DepthwiseCo (None, 4, 4, 960) 8640
block_16_expand_relu[0][0]
-----

-----
block_16_depthwise_BN (BatchNor (None, 4, 4, 960) 3840
block_16_depthwise[0][0]
-----

-----
block_16_depthwise_relu (ReLU) (None, 4, 4, 960) 0
block_16_depthwise_BN[0][0]
-----

-----
block_16_project (Conv2D) (None, 4, 4, 320) 307200
block_16_depthwise_relu[0][0]
-----

-----
block_16_project_BN (BatchNorma (None, 4, 4, 320) 1280
block_16_project[0][0]

```

```

-----
-----
Conv_1 (Conv2D) (None, 4, 4, 1280) 409600
block_16_project_BN[0][0]
-----
-----
Conv_1_bn (BatchNormalization) (None, 4, 4, 1280) 5120 Conv_1[0][0]
-----
-----
out_relu (ReLU) (None, 4, 4, 1280) 0 Conv_1_bn[0][0]
=====
Total params: 2,257,984
Trainable params: 0
Non-trainable params: 2,257,984
-----
-----

```

```

[90]: model = Sequential()

model.add(conv_base)

model.add(Flatten())

model.add(Dense(128, activation='relu'))

model.add(Dense(1, activation='sigmoid'))

```

```

[92]: model.summary()

```

```

Model: "sequential_4"

```

```

-----
Layer (type)                Output Shape                Param #
=====
mobilenetv2_1.00_128 (Functi (None, 4, 4, 1280)         2257984
-----
flatten_6 (Flatten)         (None, 20480)              0
-----
dense_12 (Dense)            (None, 128)                2621568
-----
dense_13 (Dense)            (None, 1)                  129
=====
Total params: 4,879,681
Trainable params: 2,621,697
Non-trainable params: 2,257,984
-----

```

```
[95]: model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

```
[96]: import time

start = time.perf_counter()

perf = model.fit_generator(train_generator, epochs=15,
    ↳ callbacks=[early_stop], validation_data=validation_generator)

elapsed = time.perf_counter() - start

print('Elapsed {}'.format(elapsed/60))
```

```
/usr/local/lib/python3.6/dist-packages/tensorflow/python/keras/engine/training.py:1844: UserWarning:
`Model.fit_generator` is deprecated and will be removed in a future version.
Please use `Model.fit`, which supports generators.
```

```
warnings.warn("`Model.fit_generator` is deprecated and "
```

```
Epoch 1/15
```

```
271/293 [=====>...] - ETA: 5s - loss: 0.2902 - accuracy:
0.9396
```

```
/usr/local/lib/python3.6/dist-packages/PIL/TiffImagePlugin.py:770: UserWarning:
Possibly corrupt EXIF data. Expecting to read 32 bytes but only got 0. Skipping
tag 270
```

```
" Skipping tag %s" % (size, len(data), tag)
```

```
/usr/local/lib/python3.6/dist-packages/PIL/TiffImagePlugin.py:770: UserWarning:
Possibly corrupt EXIF data. Expecting to read 5 bytes but only got 0. Skipping
tag 271
```

```
" Skipping tag %s" % (size, len(data), tag)
```

```
/usr/local/lib/python3.6/dist-packages/PIL/TiffImagePlugin.py:770: UserWarning:
Possibly corrupt EXIF data. Expecting to read 8 bytes but only got 0. Skipping
tag 272
```

```
" Skipping tag %s" % (size, len(data), tag)
```

```
/usr/local/lib/python3.6/dist-packages/PIL/TiffImagePlugin.py:770: UserWarning:
Possibly corrupt EXIF data. Expecting to read 8 bytes but only got 0. Skipping
tag 282
```

```
" Skipping tag %s" % (size, len(data), tag)
```

```
/usr/local/lib/python3.6/dist-packages/PIL/TiffImagePlugin.py:770: UserWarning:
Possibly corrupt EXIF data. Expecting to read 8 bytes but only got 0. Skipping
tag 283
```

```
" Skipping tag %s" % (size, len(data), tag)
```

```
/usr/local/lib/python3.6/dist-packages/PIL/TiffImagePlugin.py:770: UserWarning:
Possibly corrupt EXIF data. Expecting to read 20 bytes but only got 0. Skipping
tag 306
```

```
" Skipping tag %s" % (size, len(data), tag)
```

```
/usr/local/lib/python3.6/dist-packages/PIL/TiffImagePlugin.py:770: UserWarning:
```

Possibly corrupt EXIF data. Expecting to read 48 bytes but only got 0. Skipping tag 532

```
" Skipping tag %s" % (size, len(data), tag)
/usr/local/lib/python3.6/dist-packages/PIL/TiffImagePlugin.py:788: UserWarning:
Corrupt EXIF data. Expecting to read 2 bytes but only got 0.
warnings.warn(str(msg))
```

```
293/293 [=====] - 93s 309ms/step - loss: 0.2795 -
accuracy: 0.9410 - val_loss: 0.0867 - val_accuracy: 0.9635
```

Epoch 2/15

```
293/293 [=====] - 89s 303ms/step - loss: 0.0524 -
accuracy: 0.9800 - val_loss: 0.0967 - val_accuracy: 0.9651
```

Epoch 3/15

```
293/293 [=====] - 89s 303ms/step - loss: 0.0266 -
accuracy: 0.9907 - val_loss: 0.1207 - val_accuracy: 0.9610
```

Epoch 4/15

```
293/293 [=====] - 88s 302ms/step - loss: 0.0168 -
accuracy: 0.9939 - val_loss: 0.1493 - val_accuracy: 0.9619
```

Elapsed 5.979793827750003

```
[97]: new_loss, new_acc = model.evaluate_generator(validation_generator)
```

```
/usr/local/lib/python3.6/dist-
packages/tensorflow/python/keras/engine/training.py:1877: UserWarning:
`Model.evaluate_generator` is deprecated and will be removed in a future
version. Please use `Model.evaluate`, which supports generators.
warnings.warn("`Model.evaluate_generator` is deprecated and "
```

```
[98]: print('Validation Accuracy - {:.2f}%'.format(new_acc*100))
print('validation loss - {:.2f}'.format(new_loss))
```

```
Validation Accuracy - 96.19%
validation loss - 0.15
```

```
[99]: model_history = pd.DataFrame(perf.history)
```

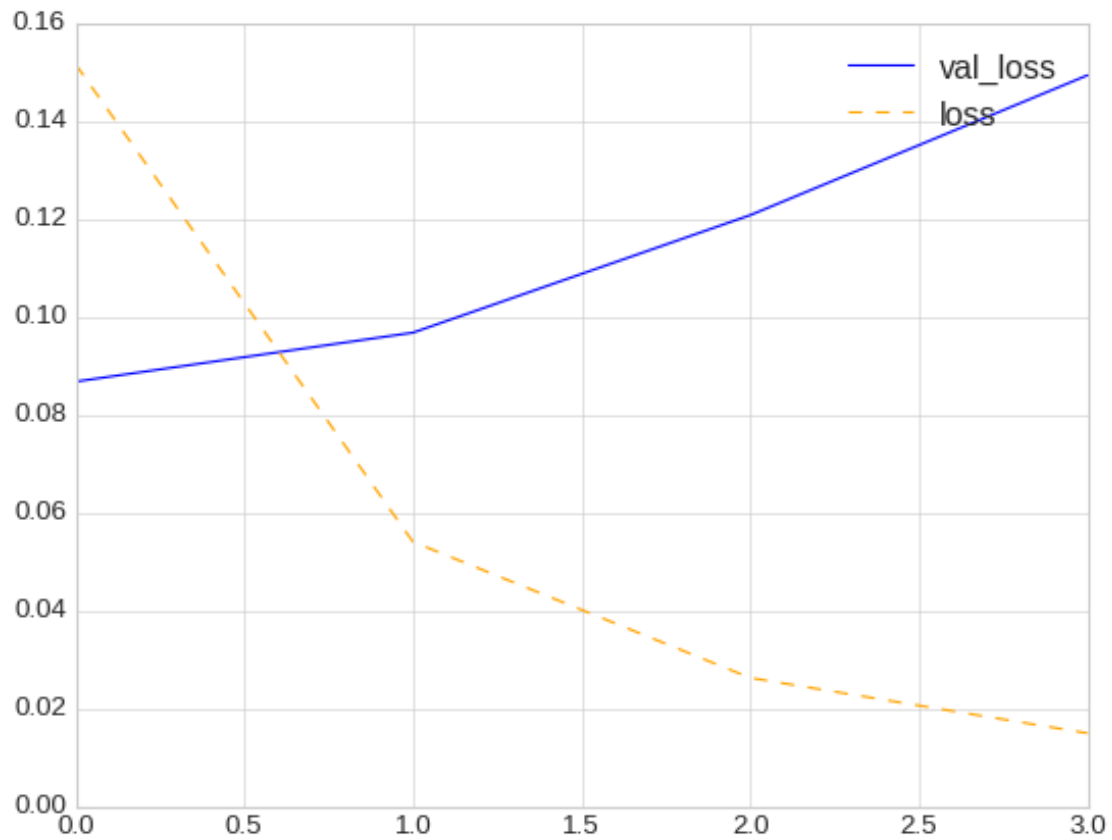
```
[100]: model_history
```

```
[100]:
```

	loss	accuracy	val_loss	val_accuracy
0	0.151495	0.958011	0.086692	0.963526
1	0.053979	0.979619	0.096691	0.965126
2	0.026196	0.990236	0.120692	0.960966
3	0.014898	0.994505	0.149257	0.961926


```
[101]: plt.style.use('seaborn-whitegrid')
plt.plot(model_history['val_loss'], ls='-', color='blue', label='val_loss')
plt.plot(model_history['loss'], ls='--', color='orange', label='loss' )
plt.legend()
```

[101]: <matplotlib.legend.Legend at 0x7f22163a81d0>



```
[103]: plt.style.use('seaborn-whitegrid')
plt.plot(model_history['val_accuracy'],ls='-',color='blue',label='val_accuracy')
plt.plot(model_history['accuracy'],ls='--',color='orange',label='accuracy' )
plt.legend(loc='upper left')
```

[103]: <matplotlib.legend.Legend at 0x7f22149be6d8>



```
[104]: pred_probs = model.predict_generator(validation_generator)
       preds = pred_probs > 0.5
```

```
/usr/local/lib/python3.6/dist-packages/tensorflow/python/keras/engine/training.py:1905: UserWarning:
`Model.predict_generator` is deprecated and will be removed in a future version.
Please use `Model.predict`, which supports generators.
  warnings.warn("`Model.predict_generator` is deprecated and "
```

```
[105]: print(classification_report(preds, validation_generator.classes))
```

	precision	recall	f1-score	support
False	0.98	0.95	0.96	3236
True	0.94	0.98	0.96	3015
accuracy			0.96	6251
macro avg	0.96	0.96	0.96	6251
weighted avg	0.96	0.96	0.96	6251


```
[106]: plt.figure(figsize=(10,7))
plt.suptitle('Confusion Matrix')
confmat = pd.DataFrame(confusion_matrix(preds, validation_generator.classes),
    ↪ columns=validation_generator.class_indices.keys())
confmat.index = validation_generator.class_indices.keys()
sns.heatmap(confmat, annot=True, fmt='d', cmap='Blues')
```

```
[106]: <matplotlib.axes._subplots.AxesSubplot at 0x7f22163e8d68>
```



1 Inference-2

With **MobileNetV2** + custom output layer, the model achieved 99% on training set and 96% on the Validation Set, this difference indicates that the model is overfitting on the training set, but this is not prominent to finetune further generally. But to create a state of the art cats & dogs classifier we can further fine-tune it to reduce that 3% of variance with techniques like Dropout and L2 Regularisation.

```
[107]: model.save('Cats&Dogs-MovbileNetv2.h5')
```

1.1 Miscellaneous

```
[108]: import os
path = './samples/'
overview_path = './samples/overview.txt'
eval_path = './samples/evaluate.txt'

if os.path.exists(path):

    print('samples dir, exists..checking for dictionaries existence..')

    if os.path.exists(overview_path) and os.path.exists(eval_path):
        print('Data exists. no need of overwritting.')
    else:
        print("overview and eval doesn't exist, proceed to step-2")

else:
    print("samples/ dir is non-existent, Establishing one..")
    os.mkdir(path) # samples directory
```

samples/ dir is non-existent, Establishing one..


```
[110]: x_train, y_train = next(train_generator) # one batch
x_test, y_test = next(validation_generator) # one batch
```

```
[136]: classes = np.array(list(validation_generator.class_indices))
x = np.array((x_train[0],x_train[7], x_train[10]))
y = np.array((y_train[0],y_train[7], y_train[10]))
```

```
[147]: y = ['CAT', 'DOG', 'CAT']
```

```
[148]: plt.imshow(x[2])
plt.title(y[2])
```

```
[148]: Text(0.5, 1.0, 'CAT')
```



[155]:

```
[155]: array([0., 0., 0., 0., 0., 0., 0., 1., 0., 1., 0., 1., 0., 0., 0., 0., 0., 0.,
          1., 1., 1., 1., 0., 0., 0., 0., 0., 0., 1., 0., 0., 1., 1., 1., 1.,
          1., 0., 0., 0., 0., 0., 1., 1., 1., 1., 0., 1., 0., 1., 0., 1., 0.,
          1., 0., 0., 1., 0., 1., 0., 0., 1., 0., 1., 0., 1.], dtype=float32)
```

[157]:

```
overview_dict = {}
eval_dict = {}

# fill the following -
# for overview
#string
kind = 'Image Data'
#tuple
```

```

dimensions = x_train.shape
#labels : str(list of unique target values)
targets = list(validation_generator.class_indices.values())
#nd.array
data = x
#nd.array or class_names
labels = y

vars0 = ['kind', 'dimensions', 'targets', 'data', 'labels']

# filling overview_dict
for x in vars0:
    try:
        overview_dict[x] = eval(x)
    except:
        overview_dict[x] = x

# evaluate_dict

eval_dict = {'test_cases' : x_train, 'true': y_train, 'class_names':
    ↪ ['CAT', 'DOG'] , 'model': '/model.h5'}

```

```

[159]: import pickle
# dump 1
with open(overview_path, 'wb') as f:
    pickle.dump(overview_dict, f)

# dump 2
with open(eval_path, 'wb') as f:
    pickle.dump(eval_dict, f)

```

```

[160]: model.summary()

```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
mobilenetv2_1.00_128 (Func	(None, 4, 4, 1280)	2257984
flatten_6 (Flatten)	(None, 20480)	0
dense_12 (Dense)	(None, 128)	2621568
dense_13 (Dense)	(None, 1)	129
Total params: 4,879,681		
Trainable params: 2,621,697		

Non-trainable params: 2,257,984

```
[161]: s = '''
Model: "sequential_4"

-----
Layer (type)                 Output Shape              Param #
-----
mobilenetv2_1.00_128 (Functi (None, 4, 4, 1280)        2257984
-----
flatten_6 (Flatten)         (None, 20480)             0
-----
dense_12 (Dense)             (None, 128)               2621568
-----
dense_13 (Dense)             (None, 1)                 129
=====
Total params: 4,879,681
Trainable params: 2,621,697
Non-trainable params: 2,257,984
-----

'''
```

```
[162]: report = '''

           precision    recall  f1-score   support

 False      0.98        0.95        0.96        3236
  True      0.94        0.98        0.96        3015

 accuracy              0.96        6251
 macro avg           0.96        0.96        0.96        6251
weighted avg           0.96        0.96        0.96        6251

'''
```

```
[ ]:
```

```
[166]: synopsis = '''

For this problem two Implementations were used with Two inferences at the end_
↳ of each implementation

exe-1 = ConvNet from Scratch
```

- * A ConvNet was built from scratch with 3 Conv2D, 3 MaxPool, 1 Flatten and 2 Dense Layers with Output function as Sigmoid (you can find the implementation in the ipynb-pdf).
- * This ConvNet was set for 15 epochs with EarlyStopping Callback of patience 3. The Elapsed Training time was 12.18 Mins on Cloud GPU instance of only 8 epochs of the stipulated 15.
- * Attained Training Accuracy - 95.60% and Validation Accuracy - 83.75% with approx. 8% difference, This network was overfitting with an average f1-score of 84%.
- * This Network had the classic case of misclassifying cats as dogs.

```
exe-2 = MobileNetv2 ( Frozen )
```

- * Convbase of MobileNetV2 with ImageNet weights and Frozen layers was constructed.
- * Custom Output layers of 1 Flatten, 2 Dense was attached at the end.
- * Same as the previous implementation the model was set to 15 epochs with same EarlyStopping Setup. Elapsed Training time is 5 Mins with 3 Epochs of Training.
- * This Time the Model Performed well enough with the pre-trained ImageNet Weights. The misclassification problem from the first case was smoothened.
- * With **MobileNetV2** + custom output layer, the model achieved 99% on training set and 96% on the Validation Set, this difference indicates that the model is overfitting on the training set, but this is not prominent to fine tune further generally. But to create a state of the art cats & dogs classifier we can further fine-tune it to reduce that 3% of variance with techniques like Dropout and L2 Regularisation.

```
'''
```

```
[167]: print(synopsis)
```

For this problem two Implementations were used with Two inferences at the end of each implementation

```
exe-1 = ConvNet from Scratch
```

- * A ConvNet was built from scratch with 3 Conv2D, 3 MaxPool, 1 Flatten and 2 Dense Layers with Output function as Sigmoid (you can find the implementation in the ipynb-pdf).
- * This ConvNet was set for 15 epochs with EarlyStopping Callback of patience 3. The Elapsed Training time was 12.18 Mins on Cloud GPU instance of only 8 epochs of the stipulated 15.

* Attained Training Accuracy - 95.60% and Validation Accuracy - 83.75% with approx. 8% difference, This network was overfitting with an average f1-score of 84%.

* This Network had the classic case of misclassifying cats as dogs.

exe-2 = MobileNetv2 (Frozen)

* Convbase of MobileNetV2 with ImageNet weights and Frozen layers was constructed.

* Custom Output layers of 1 Flatten, 2 Dense was attached at the end.

* Same as the previous implementation the model was set to 15 epochs with same EarlyStopping Setup. Elapsed Training time is 5 Mins with 3 Epochs of Training.

* This Time the Model Performed well enough with the pre-trained ImageNet Weights. The misclassification problem from the first case was smoothened.

* With **MobileNetV2** + custom output layer, the model achieved 99% on training set and 96% on the Validation Set, this difference indicates that the model is overfitting on the training set, but this is not prominent to fine tune further generally. But to create a state of the art cats & dogs classifier we can futhur fine-tune it to reduce that 3% of variance with techniques like Dropout and L2 Regularisation.

```
[170]: desc = '''A simple Cats and Dogs Image Dataset with 24994 samples of both cats_
↳and dog pictures combined.'''
project_name = 'Cats & Dogs Classifier'
framework = 'Keras'
prediction_type = 'Classification of 2 Targets'
network_type = 'MobileNetV2'
architecture = s
layers = '19 Residual Units + 3 Custom Output Layers'
hidden_units = 2
activations = "['relu', 'sigmoid']"
epochs = "Set of 15, Trained for 3 ( EarlyStopping ) "
metrics = "Accuracy"
loss = "Binary Cross-Entropy"
optimiser = 'Adam'
learning_rate = 0.001
batch_size = 64
train_performance = '99.45%'
test_performance = '96.14%'
classification_report = report
elapsed = "5Min, runtime : Colab Cloud GPU"
summary = synopsis
ipynb = './Projects/Transfer-Learning/Dogs&Cats/Dogs&Cats-classifier.pdf'
plots = './Projects/Transfer-Learning/Dogs&Cats/Plots'
```

```
[184]: p = {}
```

```
[194]: var = ['desc', 'project_name',  
            ↪ 'framework', 'prediction_type', 'network_type', 'architecture', 'layers', 'hidden_units', 'activation'  
  
param = {}  
for val in var:  
  
    try:  
        param[val] = eval(val)  
  
    except:  
        param[val] = val  
  
    # check if anything is missing
```

```
[195]: param
```

```
[195]: {'activations': '["relu', 'sigmoid']",
        'architecture': '\nModel: "sequential_4"\n
        -----\n
        \nLayer (type)                                     Output Shape
Param # \n===== \n
obilenetv2_1.00_128 (Function (None, 4, 4, 1280)         2257984
\n
\nflatten_6
(Flatten)          (None, 20480)                0
\n
\nndense_12
(Dense)            (None, 128)                 2621568
\n
\nndense_13
(Dense)            (None, 1)                   129
\n
\n===== \nTotal
params: 4,879,681\nTrainable params: 2,621,697\nNon-trainable params: 2,257,984\n
\n----- \n\n\n\n',
        'batch_size': 64,
        'classification_report': '\n
        \n
        precision      recall
f1-score  support\n\n      False    0.98    0.95    0.96    3236\n
True      0.94    0.98    0.96    3015\n\n
        accuracy
0.96    6251\n
        macro avg    0.96    0.96    0.96    6251\n
        weighted
avg      0.96    0.96    0.96    6251\n\n\n',
        'desc': 'A simple Cats and Dogs Image Dataset with 24994 samples of both cats
        and dog pictures combined.',
        'elapsed': '5Min, runtime : Colab Cloud GPU',
        'epochs': 'Set of 15, Trained for 3 ( EarlyStopping ) ',
        'framework': 'Keras',
        'hidden_units': 2,
        'ipynb': './Projects/Transfer-Learning/Dogs&Cats/Dogs&Cats-classifier.pdf',
        'layers': '19 Residual Units + 3 Custom Output Layers',
        'learning_rate': 0.001,
```



```

'loss': 'Binary Cross-Entropy',
'metrics': 'Accuracy',
'network_type': 'MobileNetV2',
'optimiser': 'Adam',
'plots': './Projects/Transfer-Learning/Dogs&Cats/Plots',
'prediction_type': 'Classification of 2 Targets',
'project_name': 'Cats & Dogs Classifier',
'summary': '\n\nFor this problem two Implementations were used with Two
inferences at the end of each implementation\n\nexe-1 = ConvNet from
Scratch\n\n* A ConvNet was built from scratch with 3 Conv2D, 3 MaxPool, 1
Flatten and 2 Dense Layers with Output function as Sigmoid ( you can find the
implementation in the ipynb-pdf ). \n* This ConvNet was set for 15 epochs with
EarlyStopping Callback of patience 3. The Elapsed Training time was 12.18 Mins
on Cloud GPU instance of only 8 epochs of the stipulated 15.\n* Attained
Training Accuracy - 95.60% and Validation Accuracy - 83.75% with approx. 8%
difference, This network was overfitting with an average f1-score of 84%.\n*
This Network had the classic case of misclassifying cats as dogs.\n\nexe-2 =
MobileNetv2 ( Frozen )\n\n* Convbase of MobileNetV2 with ImageNet weights and
Frozen layers was constructed.\n* Custom Output layers of 1 Flatten, 2 Dense was
attached at the end.\n* Same as the previous implementation the model was set to
15 epochs with same EarlyStopping Setup. Elapsed Training time is 5 Mins with 3
Epochs of Training.\n* This Time the Model Performed well enough with the pre-
trained ImageNet Weights. The misclassification problem from the first case was
smoothened. \n* With **MobileNetV2** + custom output layer, the model achieved
99% on training set and 96% on the Validation Set, this difference indicates
that the model is overfitting on the training set, but this is not prominent to
fine tune further generally. But to create a state of the art cats & dogs
classifier we can futhur fine-tune it to reduce that 3% of variance with
techniques like Dropout and L2 Regularisation.\n\n\n',
'test_performance': '96.14%',
'train_performance': '99.45%'}

```

```

[196]: import pickle
file = open("artefacts.txt", "wb")
dictionary = param
pickle.dump(dictionary, file)
file.close()

```

```
[ ]:
```